
Reliable and Invalid: Anatomy of a Gamed Training Task

Anonymous Author(s)

Affiliation

Address

email

Abstract

Verifiable rewards are trusted because they are deterministic: a program either awards the points or it does not. We post-train Qwen2.5-7B-Instruct with GRPO on NYT Connections under a reward with a cheap structural component (well-formed output) and an expensive semantic component (correct groupings). On a strictly chronological test set of 162 puzzles (648 group slots), scored in one serving session, capability runs $26 \rightarrow 52 \rightarrow 80 \rightarrow 4$ slots across the untrained Instruct model, supervised fine-tuning, GRPO step 50, and the GRPO endpoint: paired stage gains of $+0.161$ and $+0.173$ on the 0–4 scale, then a fall below the untrained model in all three seeds — 4, 7 and 11 of 648 slots. The *training* reward reads the destruction as mastery: it reaches its 1.6 ceiling by step 125 with zero across-rollout variance and stays there. Scoring the endpoint offline on its own training puzzles — 21 of 3,228 slots — rules out memorization and exposes the cause: a one-line bug in our GRPO data loader presented each board’s sixteen words in answer-key order, so the training task, as presented, was solvable by copying the words four at a time. The collapsed policy does exactly that — 91.4% of the groups it emits on training boards, and 92.4% on test boards, are consecutive quadruples of its prompt; a copy of a shuffled board scores at or below the chance floor, which is why the endpoint’s score collapses, and its small margin above that floor comes from the residual non-copy responses. Entropy collapse, cross-seed convergence, and the reward ceiling are signatures of one deterministic copy rule. The reward function was correct on every string it scored; the task presentation was gameable, and every in-sample signal read the gaming as mastery. The same reward held out, and three blind LLM judges, catch the reversal; four lines of string matching identify the rule; and rerunning the recipe with the one-line fix, the same reward trains the best policies in this study — 148, 135 and 129 of 648 slots across three seeds, each about five times the untrained model, with no collapse, no copying, and every in-sample signal reading honestly. Supervised fine-tuning on the leaked prompts alone converges to a perfect copier: the gameable object was the task, and it gamed every learner we gave it. Determinism certifies the scorer, not the construct: hold the construct out, not just the reward.

1 Introduction

Post-training against verifiable rewards is attractive for a reason that sounds decisive: the reward is a program. It returns exactly what it was written to return, and the argument that this buys immunity from reward hacking¹ rests on that reliability. This paper is about the gap reliability does not close: the gap between the reward and the construct it stands in for. We train on NYT Connections — sixteen

¹We follow the common usage *reward hacking* throughout; Skalse et al. [2022], whose formalization we adopt, publish it under the name *reward gaming*.

words, partitioned into four groups of four — chosen because well-formedness and correctness are unusually separable: a response can be a perfectly valid partition and completely wrong. Our reward pays for both; two of its five components can be checked by reading the string, three require the answer key.

For fifty steps it works. In one serving session, held-out groups recovered run $26 \rightarrow 52 \rightarrow 80$ across the untrained Instruct model, the supervised checkpoint GRPO is initialized from, and step 50 — paired stage gains of $+0.161$ [$+0.056, +0.272$] and $+0.173$ [$+0.074, +0.284$], the reinforcement-learning stage contributing as much as supervised fine-tuning in the one session pairing all four arms (step 50 decodes 77–80 and SFT 52–56 across sessions: the two stage gains are indistinguishable, not ordered; §3). Over the following 353 steps the same run falls to 4 slots, below the untrained model and to within $3\times$ of the 1.4-slot floor of a uniformly drawn valid partition, while the invalid-output rate falls by an order of magnitude and the training reward reaches its 1.6 maximum and averages 1.5814 thereafter, with zero variance across the $K = 8$ rollouts on $\sim 95\%$ of prompts. A practitioner watching that curve saw a solved task. The curve was not lying about the reward; it was lying about the task. A control run scoring the endpoint on its own training puzzles returned 21 of 3,228 slots — not memorization — and led to the mechanism: our GRPO data loader, unlike our SFT and evaluation code, presented each board’s words in answer-key order, and the converged policy simply copies the prompt four words at a time (91–92% of emitted groups; §7). The failure is not in the reward function, which is correct; it is in what the training distribution let the reward certify, and in the practice of reading the certificate in sample. The setting matters: in RL post-training the training distribution is regenerated by the trainer at every step, so a leak like ours lives in code rather than in a dataset — where none of the auditing reflexes the shortcut-learning literature built will look for it.

Contributions. A fully verifiable case study in which GRPO improves and then destroys the target construct, ending below the untrained model across three seeds (§3); a complete post-hoc diagnosis: an answer-ordered prompt bug made the training task gameable, and the converged policy is a prompt-order copier on 91–92% of its emitted groups (§7); an account in which entropy collapse, the reward ceiling, and cross-seed convergence are signatures of that one rule (§5), surviving a normalization ablation (§6); the lesson demonstrated end to end: every in-sample signal read gaming as mastery; the reward held out, and three blind judges, caught it (§7–§9); and the diagnosis closed by intervention: with the fix, the recipe trains to 148/135/129 across three seeds with no collapse, while SFT on the leaked prompts alone converges to a perfect copier — the attractor is the task’s, not the algorithm’s (§8).

2 Task, pipeline, and the reward’s asymmetry

Task and data. 1,078 dated NYT Connections puzzles (1,077 after validation), split strictly chronologically: 807 train (2023-06-12 to 2025-08-27), 108 validation, 162 test (2025-12-15 to 2026-05-29); every test puzzle postdates every training puzzle. The game has prior use as an LLM benchmark [Todd et al., 2024]. Date regressions of groups-correct across the test window are flat for both base and the trained policy (values in the repository): contamination controlled, though the intervals are wide and only large drift would show.

Pipeline. Qwen2.5-7B-Instruct (and 1.5B for the scale contrast), rank-16 LoRA ($\alpha = 32$, dropout 0.05, all linear modules) on a 4-bit base. SFT for 3 epochs on the training split (learning rate 10^{-4} , batch size 2, gradient accumulation 8), then GRPO [Shao et al., 2024] from the SFT checkpoint: $K = 8$ sampled completions per puzzle at temperature 0.9, group-relative advantages, $\beta = 0.001$ against the SFT reference, learning rate 5×10^{-6} , batch size 2 with gradient accumulation 8, one epoch (403 optimizer steps) at 7B.

The leak, disclosed up front. SFT data and every evaluation in this paper present each board’s sixteen words deterministically shuffled (seeded by puzzle id), precisely so group order is not a giveaway. Our GRPO loader did not: a one-line bug built training prompts from the answer key, so during reinforcement learning — only there — the words arrived grouped in answer order: every GRPO prompt was solvable by a positional copy rule worth the full 1.6. Discovered post hoc (§7); disclosed here because nothing in the training-time sections can be read without it. Held-out measurements are unaffected: evaluation prompts were always shuffled. §8 reruns the recipe with the fix.

89 **The reward.** Table 1 is the instrument every later section reads.

Table 1: The reward decomposition. Two components are verifiable from the emitted string alone; three require the answer key. The policy is trained on the sum.

Component	Value	Cost to check
Parses as four groups of four	+0.1	cheap — read the string
Correct groups	$+1.0 \times (\text{correct}/4)$	needs the answer key
All four groups correct	+0.5	needs the answer key
Group with exactly 3 correct words	+0.05	needs the answer key
Invalid output	−0.1	cheap — read the string

90 Maximum 1.6. The model can raise its score by getting better at the puzzle or by getting better at
 91 emitting well-formed output; only one is the construct. That asymmetry — cheap and expensive
 92 components, separately reportable — did not cause the failure: the copy rule earned the expensive
 93 components too (§7). Its role is diagnostic: the failure is invisible in the aggregate, obvious in the
 94 decomposition (§10).

95 3 The arc on held-out data

96 Table 2 reports all four arms of the reported 7B run on the 162-puzzle test set, greedy, one vLLM
 97 session so every comparison is within-session; the supervised checkpoint is included because it, not
 98 the untrained model, is what GRPO initialized from.

Table 2: 7B on the test split (162 puzzles, 648 group slots; all four arms in one serving session; greedy). Means on the 0–4 scale are exact ratios of the count column. Bootstrap intervals are percentile over puzzles, $B = 1,000$ resamples, seed 0. This is the session §7 also scores on the training split.

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward	Solves
Instruct (untrained)	0.1605 [0.105, 0.222]	26	6.8%	0.165	0/162
SFT	0.3210 [0.216, 0.432]	52	22.2%	0.191	2/162
GRPO, step 50	0.4938 [0.383, 0.623]	80	16.0%	0.252	2/162
GRPO, step 403	0.0247 [0.000, 0.056]	4	0.6%	0.125	0/162

99 Two reference points frame the table. A uniformly drawn valid partition recovers $4 \times$
 100 $5775/2,627,625 = 0.0088$ groups per puzzle, or **1.4 slots** across this split; the step-403 endpoint sits
 101 at 2.8 times that floor — and §7 shows why: its copied responses score at or below chance, and the
 102 excess over the floor rides on the few responses that are not copies. Second, because all four arms sit
 103 in one session, the stage accounting is paired: $26 \rightarrow 52 \rightarrow 80 \rightarrow 4$, the supervised stage contributing
 104 $+0.1605 [+0.0556, +0.2716]$ and the reinforcement-learning stage $+0.1728 [+0.0741, +0.2840]$ —
 105 near-equal gains, both intervals excluding zero.

106 The method works, then reverses. Both movements are paired within Table 2’s session: the climb to
 107 step 50 is the two stage gains just given, each excluding zero, and step 403 falls below the untrained
 108 model by $-0.1358 [-0.1975, -0.0802]$. An earlier three-arm session read the same two contrasts
 109 consistently (values in the repository); means in Table 2 are exact ratios of the count column. Two
 110 details locate the mechanism. Held-out mean reward reads 0.252 at step 50 and 0.125 at step 403
 111 against base’s 0.165 (the later-session step-100 point is in Figure 2), while the training reward reaches
 112 its ceiling by step 125 (§5): the same reward ranks the checkpoints correctly out of sample and fails
 113 to on the training curve — that contrast, not the reward’s determinism, is this paper’s subject. And
 114 the step-50 policy is worse than base on structure (16.0% invalid against 6.8%) but *not* worse than
 115 its actual initialization, the SFT arm at 22.2% — so what is specific to the phase after step 50 is the
 116 semantic collapse, not the structural climb, which was already more than half complete before the
 117 peak (§5).

118 **How wide is the peak?** In a separate session, step 100 evaluates at 76/648 to a step-50 re-serve’s
 119 80/648, paired $-0.0247 [-0.148, +0.093]$: uninformative rather than equivalence (the interval spans

the whole step-50 effect; no margin was pre-specified). The ablation trajectory suggests decline by step 100; all both sessions agree on is that the collapse falls in 100–150 (§5).

Reproducibility across serving sessions. Base and the final checkpoint reproduce per-puzzle *exactly* across vLLM sessions (a deterministic copier and a converged base leave no argmax ties to resolve); step 50 does not (77–80 of 648), which is why every step-50 number carries its serving session, and SFT decodes 52–56 (Appendix C). No conclusion turns on the differences.

Robustness of the endpoint. Three independent 7B seeds land at 4, 7 and 11 recovered slots of 648 (0.045 ± 0.022 on the 0–4 scale) — all below base’s 26. Their LoRA update directions agree pairwise (cosine +0.67 to +0.69) — three runs finding the same attractor (§7); seeds share an initialization and dataset, so no random-direction null applies. Best-of-16 sampling does not rescue the converged policy: a copier has nothing new to sample (single-seed, illustrative contrasts in Appendix C).

Conditioning on well-formed output widens the gap rather than explaining it. Is the effect an artifact of validity, since invalid generations score zero groups and validity is what GRPO improved? Restricting to structurally valid generations at 7B, each arm’s conditional computed within a single serving session (the SFT arm’s is its 56-slot session): base **0.172** ($n = 151$ of 162), SFT **0.444** ($n = 126$), GRPO **0.025** ($n = 161$). This is a rescaling and we label it as one (each conditional is the unconditional mean divided by the validity rate, so the gap widens automatically when the low-scoring arm is the more valid, as here); it licenses only the negative claim we need: the deficit survives with formatting held fixed. Among answers that parse, the converged policy recovers 4 group slots where the untrained model recovers 26.

The loss shows no stratum gradient we can resolve. Relative loss from the untrained model to the endpoint runs 67–100% across all four stratum labels (Appendix D) — registered pre-diagnosis as a failed prediction of selective inference; post-diagnosis, the copy rule’s expected signature, since a positional rule is blind to stratum. With 21–69 puzzles per stratum and base rates as low as one slot, the design cannot separate uniform loss from a gradient: consistent with stratum-blindness, not evidence for it.

4 The scale contrast at 1.5B

At 1.5B the instrument reads nothing. The untrained model recovers one group slot of 648, GRPO one, and SFT — the best arm — two, against Section 3’s 1.4-slot chance floor. No arm is distinguishable from that floor (the exact intervals of Appendix F contain it), so no claim about the construct is available at this scale. What is visible is the reward driving its verifiable component: invalid output 32.1% $\rightarrow$ 2.5%, mean reward 0.049 $\rightarrow$ 0.113 (Appendix F). We report it as an observation: epoch-asymmetric runs, no 1.5B trajectory — and “untouched at 1.5B” would be an endpoint claim, which this paper shows cannot settle.

5 Trajectory of the collapse: entropy, KL, and the training signal

Figure 1 scores checkpoints of the reported run on the $n = 100$ temperature-0.9 *entropy sweep* (denominators of 400; full per-checkpoint values in Appendix B); the ablation’s *checkpoint sweep* (§6) is $n = 108$, greedy, /432 — different measurements, never mixed.

The semantic peak has a location; the structural climb does not. Semantic rises 0.0325 $\rightarrow$ 0.0950 from the SFT initialization to step 50 (13/400 $\rightarrow$ 38/400 slots), falls to 0.0075 by step 150, and holds at 0.0100 from 200 on. Structural (valid-partition fraction, not shown) climbs 0.37 $\rightarrow$ 0.68 $\rightarrow$ 0.64 $\rightarrow$ 0.95–0.96 across init/50/100/150 on — one reversal, at 100 — with 0.31 of its 0.59 total gain banked by step 50, concurrent with the semantic peak.

The trajectory is not a trade with a turning point: a broadly monotone structural climb (one reversal, at 100) with a semantic hump on top, and only the hump has a location. One gap qualifies the shape: no checkpoint exists between steps 100 and 150 — 81% of the KL displacement (8.12 $\rightarrow$ 46.42 nats) — so the descent is interpolated, not measured.




Figure 1: The mechanism, on the $n = 100$ temperature-0.9 entropy sweep (reported 7B run, seed 0). **Left:** policy entropy against optimizer step — a 30-fold collapse, with the cliff between steps 100 and 150. **Right:** semantic score against KL from the SFT initialization: the peak arrives by 2.70 of an eventual 47.05 nats (5.7%; step 50 is the earliest checkpoint retained), the remainder undoes it. The untrained model is an off-trajectory reference at 42.60 nats: KL measures displacement, not damage (§5).




Figure 2: The same reward, in sample and held out (reported 7B run, seed 0). Training reward (gray; sampled $T=0.9$, $K=8$; TRL 5-step averages; measured on the answer-ordered prompts of §2) first touches its 1.6 ceiling at step 125, averaging 1.5814 thereafter. Held-out mean reward (162 test puzzles; markers; greedy; the control session of Table 3; right axis, note the scale) peaks at 0.252 at step 50 and ends at 0.125, below the untrained model’s 0.165 — the dashed span between the anchors is unmeasured; the hollow step-100 point is a later session and is not joined.

167 **Entropy collapse is the rule’s signature.** Policy entropy falls 13.4-fold between steps 50 and
168 150, the window in which the semantic score falls 12.7-fold; by step 150 the policy is effectively
169 deterministic. With §7’s diagnosis in hand the reading is direct: a positional copy rule is deterministic,
170 blind to the board’s content, and reliably well-formed: entropy collapse and semantic destruction read
171 as two views of one convergence — a reading made after the diagnosis, and tested by intervention:
172 remove the leak and entropy never collapses (§8) — which establishes the leak’s necessity, not this
173 finer reading; the leaked-SFT cell of §8 adds its sufficiency without RL. The structural climb is not
174 — over half of it ($0.37 \rightarrow 0.68$ by step 50) predates measurable copying (0.3% at step 50 by the
175 strict all-four-quadruples rule of Appendix D; a graded partial-adherence statistic could date copying
176 earlier, which is why we write *measurable*): format gains the rule later locks in. What the collapse
177 removes is within-prompt variation — nothing new to sample (the pass@16 result), no group-relative
178 learning signal (the zero-advantage fraction). And what survives on held-out boards is measured
179 twice over: the cheap structural term ($0.1(1 - 2 \cdot \text{invalid})$) is 52%/27%/79% of held-out reward across
180 base/step 50/endpoint, and the copy rule is visible in 90.7% of the endpoint’s test generations (§7).

181 **KL is the budget view, and it is a path length, not a damage scale.** The step-50 checkpoint
182 sits 2.70 nats from the SFT initialization and is, jointly with step 100 (§3), the best policy this run
183 produced on held-out data. The run spent 47.05: 80 slots at 5.7% of the displacement, 4 at all of it.
184 Two qualifications keep this a path description, not a scale: step 50 is the earliest checkpoint retained,
185 so “at 5.7%” means *at or before* it; and KL measures displacement, not damage — the untrained
186 model sits 42.60 nats out on the same estimator while scoring 0.0125 on this sweep (Appendix B), so
187 base is off this trajectory and 5.7% dates the peak within this run only. The sweep also compresses
188 base’s margin over the endpoint (5/400 vs. 4/400 here; 26 vs. 4 on the greedy test instrument):
189 temperature-0.9 sampling costs a high-entropy policy far more than a collapsed one.

The training signal, read from the trainer’s own logs. Figure 2 shows the paper’s central contrast on one axis: the training reward first touches the reward function’s ceiling of 1.6000 at step 125 and averages 1.5814 over the logged points from then on, with excursions that thin as training proceeds: the deepest is 1.4300 at step 155, and the count of logged points at exactly 1.6000 rises from 14 of 28 in the first half to 24 of 29 in the second. We describe rather than test — the series is autocorrelated, the metric bounded, no trend statistic we computed is identifiable, and part of the rise is the instrument (below). What the figure cannot show is TRL’s own `frac_reward_zero_std` diagnostic: the fraction of prompts whose eight sampled rewards are identical — zero group-relative advantage. It rises from 0.0 at step 50 to 1.0 at step 125 and averages 0.946 over the remainder of training, so from step 125 onward the optimizer receives no learning signal on nearly every prompt. The saturation is legible under the diagnosis: on answer-ordered prompts (§2) a copy rule earns 1.6 on every board, seen or unseen — consistent with saturation beginning mid-epoch, on first-visit batches, and with the vanishing variance. Read as diagnostics, both in-sample signals bracket the collapse — the mean’s climb to ceiling (0.9931 $\rightarrow$ 1.6000) and the variance’s fall to zero both land between logged steps 100 and 125 — and neither signs it: saturation with zero spread is what mastery and degeneracy alike look like. Part of the climb is the instrument: as entropy collapses, temperature-0.9 sampling converges on greedy, so some of the rise is the sampling penalty vanishing (the sweep-vs-test gap, reversed). The remaining 278 steps of movement while $\sim 95\%$ of prompts carry zero advantage are consistent with continued sharpening of the rule on the residual prompts that still do; the weight-space counterpart of that push we have not traced.

β was numerically inert, and the bracket is not an experiment. An earlier configuration at learning rate 10^{-6} and $\beta = 0.04$ froze the policy ($KL \approx 0.001$; greedy outputs byte-identical to SFT); the reported one overshoots. The pair differs in two variables at once, so it bounds a region jointly and attributes nothing to either; and $\beta = 0.001$ was numerically inert ($\beta \cdot KL \leq 0.047$ reward units throughout), so this paper has no evidence about a KL penalty large enough to matter.

6 The ablation: the collapse is not the normalization default

We inherited `scale_rewards="group"`, verified as the trl 1.10.0 default. Reviewers of an earlier draft independently named it as a confound: near determinism, the within-group standard deviation dividing the advantage approaches zero, amplifying sampling noise into large updates — which alone might produce the collapse and its timing; Liu et al. [2025] argue against it on independent grounds.

We reran the 7B pipeline identically — same SFT warm start, learning rate, β , K , temperature, LoRA rank, epoch count, seed, and the same answer-ordered prompts of §2, so this ablation isolates the normalization, not the leak — with reward scaling disabled, the resolved value read back off the constructed trainer configuration. The signature is unchanged on the greedy $n = 108$ checkpoint sweep: validation semantic score rises to 62/432 at step 50, falls to 35/432 by step 100 and 4/432 by step 150, while structural validity climbs 0.70 to 0.96. The validation-selected peak (step 50, frozen before any test puzzle was scored) recovers 87/648 on test, $+0.377$ [$+0.247$, $+0.506$] paired against base; the final policy 7/648, below base by -0.117 [-0.179 , -0.056] paired and within $+0.019$ of the original endpoint.

The signature survives in the dynamics too: without reward scaling, entropy falls to 0.0128 by step 150 (original endpoint 0.0099) and KL plateaus near 46 nats (original 47.05), intermediate magnitudes differing; the $n = 100$ temperature-0.9 sweep (a different measurement) reads the same shape: 35/400 at 50, 17/400 at 100, 2/400 from 150 on. Hump, collapse, entropy endpoint, below-base endpoint: all survive with the normalization off. The phenomenon is not a library default.

7 The control run: not memorization — a copy rule

Training-reward saturation says the policy solves its training prompts under sampling; was that SFT-stage memorization? The control: four arms, one vLLM session, greedy, both full splits, the same shuffled prompts as every other evaluation (Table 3).

Three results. The endpoint retains nothing — 21 of 3,228 training slots, indistinguishable from its own test floor: not a train/test gap, so not memorization. SFT does partially memorize, though less than the raw numbers suggest: its $+0.191$ train-minus-test gap must be read against the untrained

Table 3: The control session: four arms, one vLLM session, greedy, both splits, shuffled evaluation prompts; percentile bootstrap, seed 0. Copy rate = emitted groups equal to consecutive quadruples of the prompt’s word order (Appendix D).

Arm	Train groups (0–4)	of 3,228	Test groups (0–4)	of 648	Copy rate (test)
Instruct (untrained)	0.1871 [0.156, 0.224]	151	0.1605 [0.105, 0.222]	26	2.6%
SFT	0.5118 [0.461, 0.568]	413	0.3210 [0.216, 0.432]	52	0.3%
GRPO, step 50	0.6803 [0.616, 0.746]	549	0.4938 [0.383, 0.623]	80	0.2%
GRPO, step 403	0.0260 [0.012, 0.041]	21	0.0247 [0.000, 0.056]	4	92.4%

model’s own +0.027 gap on the same chronological splits (split difficulty, not memorization); the adjusted excess is +0.164. Step 50’s adjusted excess is +0.160 to SFT’s +0.164 — a gap far inside the split-difference’s uncertainty, so inheritance anywhere from none to all is consistent — while it improves on SFT by the same margin on both splits (+0.169 / +0.173): the peak’s gain generalizes. And the shared session yields the paired stage split: 26 → 52 → 80 → 4, SFT +0.1605 [+0.0556, +0.2716], RL +0.1728 [+0.0741, +0.2840], endpoint −0.1358 [−0.1975, −0.0802] vs. the untrained model (the original session read −0.136 [−0.191, −0.080]).

Four lines of string matching find the mechanism. A 1.6 training reward against 21 recovered training slots forced a look at the training prompts, where the bug of §2 was found. The behavioral test counts emitted groups equal, as sets, to consecutive quadruples of the prompt’s word order (definition and full counts in Appendix D). Such quadruples are 91.4% of the groups the endpoint emits on training boards and 92.4% on test boards; 88.2% and 90.7% of its responses are pure four-quadruple copies; SFT and step 50 sit at 0.3%, below even the untrained model’s weak bias (4.9% / 2.6%). The converged policy is a prompt-order copier. On training prompts, where order was the answer, copying earns 1.6; on shuffled prompts a copy is a fixed valid partition of a random order — exactly the uniformly drawn valid partition behind §3’s 1.4-slot floor. The realized shuffles set a copier’s exact ceiling (artifact `results-analysis/sep06/empirical_floor.json`): 3 of 807 training boards and 1 of 162 test boards contain an aligned quadruple, against uniform expectations 7.1 and 1.4 — the puzzle-seeded draw is fixed, so the realized counts, not the expectations, bound a copier. The 711 pure copies recover 2 of that ceiling of 3, and §8’s perfect copier recovers exactly the 1 test slot available, on the one aligned board: copies harvest the alignments that exist and nothing else (no test attached: a puzzle’s four slots are not independent). What sits above chance is the aggregate — 21 slots against 7.1 expected; Table 3’s puzzle-level interval, [0.012, 0.041] on the 0–4 scale, excludes the chance mean of 0.0088, and test reads 4 against 1.4, same direction, interval not resolving — and it is carried almost entirely by responses that are *not* pure copies (35 non-copies recover 16 of the 21 slots; 60 partial copies, 3; 711 pure copies, 2): residual solving the rule has not displaced, not copying that works.

8 Reruns with the fix, and the leak without RL

We reran the reported recipe identically — same SFT warm start, hyperparameters, seed, trainer version, one epoch of 403 steps; only the logging cadence differs (Appendix E) — with the one-line fix applied, the constructed training set verified shuffled — by the same puzzle-seeded RNG as SFT and every evaluation — before any optimizer step. Everything this paper diagnoses disappears (Table 4).

The final checkpoint is the best policy in this study: 148 of 648 held-out slots — $5.7\times$ the untrained model — at +0.7531 [+0.6111, +0.9074] paired against base and +0.5802 [+0.4444, +0.7407] against SFT. Read as a recipe, not only a control: 7B QLoRA GRPO, one T4, about five hours; three seeds read 148/135/129 of 648, fully solving 10/7/6 of 162. And it is the *endpoint*: the validation curve climbs from 0.1065 to 0.1944 at step 50 to 0.1944 at step 300 and plateaus (all nine checkpoints in Appendix E); final vs. validation peak is +0.0247 [−0.0309, +0.0926] — there is no collapse to locate. Ten of 162 test puzzles are fully solved — the study’s first meaningful solve rate — and the copy rule is gone (1 of 648 emitted groups, zero pure copies).

Every in-sample pathology vanishes with it. Over the 160 logged training rows the reward never touches 1.6 — never 1.0 (maximum 0.705); `frac_reward_zero_std` peaks at 0.2 against the leaked run’s sustained 1.0; the trainer’s logged policy entropy never collapses (full range 0.116–0.212, a

Table 4: Above the rule: the leak-free session (four arms, one vLLM session, greedy, test split; conventions as Table 3; step 300 is the validation-selected peak, frozen before any test scoring). Below it: the sep05 session — seeds 1–2 of the repaired recipe and SFT trained on the leaked prompts — whose own base and SFT anchors reproduce 26 and 54.

Arm	Groups correct (0–4)	of 648	Invalid	Mean reward	Copy rate
Instruct (untrained)	0.1605 [0.105, 0.222]	26	6.8%	0.165	2.6%
SFT	0.3333 [0.228, 0.444]	54	22.2%	0.194	0.2%
GRPO-shuffled, step 300	0.8889 [0.741, 1.043]	144	6.2%	0.402	0.2%
GRPO-shuffled, step 403	0.9136 [0.759, 1.074]	148	3.7%	0.417	0.2%
GRPO-shuffled, seed 1	0.8333 [0.691, 0.988]	135	5.6%	0.388	0.2%
GRPO-shuffled, seed 2	0.7963 [0.654, 0.951]	129	5.6%	0.370	0.2%
SFT on leaked prompts	0.0062 [0.000, 0.019]	1	0.0%	0.118	100.0%

1.8 $\times$ spread against the leaked run’s 30-fold fall on the \$5 sweep — different instruments, same verdict); and the trainer’s logged KL stays below 0.23 throughout. Same reward function, same trainer: presented honestly, the task trains honestly, and the in-sample curve now sits near the held-out number (0.37 sampled vs. 0.417 greedy) instead of certifying a shortcut. The notebook’s prespecified reading for this outcome, committed before launch: the leak was the cause, its in-sample signatures intrinsic to nothing about GRPO on this task. Two further seeds of the repaired recipe land at 135 and 129 of 648 (+0.673 [+0.525, +0.815] and +0.636 [+0.500, +0.772] paired against base), with no collapse anywhere: the 100–150 window, checkpointed every 10 steps on these seeds, climbs smoothly in both; validation peaks sit at steps 200 and 250 against seed 0’s 300, and their test evaluations (127 and 132 of 648, in their own session) sit near the finals.

The leak, without RL. The remaining cell of the {leaked, clean} $\times$ {SFT, GRPO} design: SFT trained on the answer-ordered prompts — three epochs, the clean recipe otherwise — and evaluated on the same shuffled test split. It is a *perfect* copier: all 162 parsed responses are pure four-quadruple copies of their prompts (copy rate 1.000, past the leaked GRPO endpoint’s 0.924), recovering 1 of 648 slots, -0.154 $[-0.222, -0.093]$ below the untrained model, at 0.0% invalid. Plain supervised learning suffices: the copying attractor belongs to the leaked task, not to the algorithm that trains on it. What GRPO specifically did was carry a *clean* initialization into that attractor (0.3% copy at its SFT init, 92.4% at its end) — and not all the way: the non-copy responses, the trace of its clean start, carry its residual slots (§7).

9 The reversal is legible from outside the reward

The leaked-run test generations from base, step 50 and the final policy were scored blind (board words and completion only) by three LLM judges (gpt-5-mini, gemini-2.5-flash, claude-sonnet-4-5), 1–10 scale, 1458/1458 parsed. All three rank step 50 above base above the final policy, all nine pairwise per-puzzle differences excluding zero under a paired bootstrap (seed 0; per-judge step-50-minus-final gaps +0.60 to +3.25). That ordering rules out the obvious confound — the final policy is the *most* well-formed arm on this split (0.6% invalid against 6.8% and 16.0%), so a preference for well-formed text would have placed it first — in hindsight, they were reading the copy rule (§7). Judges err in correlated ways [Kohli, 2026]: three correlated readings of one signal — corroboration, not independent confirmation. The nine comparisons are one uncorrected family, all excluding zero by margins no correction would close.

10 Implications for evaluation practice

Hold out the construct, not just the reward. The strong held-out signal here is never the aggregate reward but the construct scored directly: groups correct separates endpoint from base 26 to 4 and resolves SFT’s doubling, 26 to 52, which the held-out reward cannot (SFT’s paired contrast does not exclude zero; the endpoint’s does, at just -0.040). The reward held out catches this reversal — the corollary — but 79% of the endpoint’s held-out reward is the cheap structural term: a slightly different weighting could have kept the aggregate above base through the collapse. Our case is easier

than the one where the advice is needed — this reward’s expensive component *is* the construct; absent an independent measure, the decomposition below is what remains.

Determinism certifies the scorer, not the task. Our reward was correct on every string it scored and the run was still gamed — through the prompt distribution; nothing in-sample prompted the four-line audit. One function, two measurements: in sample, reliable and invalid — this paper’s title; held out, what validity it had.

Report reward components separately, never the aggregate. The one recommendation we would keep, and it costs nothing: §5 reports it for our three arms (52%, 27%, 79%) — a standard reward-model diagnostic, simply not applied to deterministic rewards.

Assert prompt-path parity. The regression test that would have caught this bug asserts every prompt-constructing code path yields the identical string per example (Appendix A); “not in answer order” would not have. The training distribution is code, so this parity check — not a dataset audit — is the shortcut-learning defense that transfers.

Treat single-seed RL results as uninterpretable. It applies to results that come out well: the leaked run’s step 50 is one seed, the argmax of a nine-point validation grid selected before any test scoring (tabulated in Appendix B); the repaired three seeds are the caution heeded (§8).

11 Related work

Overoptimization, reasoning RL, and reward hacking. Gao et al. [2023] established the rise-then-fall curve for *learned* reward models as Goodhart against an approximation [Stiennon et al., 2020]; Karwowski et al. [2024] treat the curve formally for RL against proxies; Pan et al. [2022] report phase transitions shaped like the leaked 100–150 collapse (why that grid gap matters, §5). Our reward is a program that approximates nothing; the gaming ran through the prompt distribution (§2), a channel a better reward model does not touch.

Verifiable-reward RL underlies recent reasoning systems [DeepSeek-AI et al., 2025]; Yue et al. [2025] ask whether it extends capability beyond the base model’s pass@ k ceiling; our leaked run ends below both its antecedents, the repaired runs above them (§8). Shao et al. [2025] show Qwen2.5 gaining on MATH-500 under RLVR with random or incorrect rewards; part of the leaked step-50 gain may be that phenomenon. Cui et al. [2025] document entropy collapse as central to RL for reasoning; here it is symptom, not cause: leak removed, entropy never collapses (§8). Liu et al. [2025] critique group-level normalization; §6 shows independence from it.

Skalse et al. [2022] formalize the failure mode (as *reward gaming*); ours is gaming of the composed task — reward plus prompt distribution — not of the reward function, which was exact. Baker et al. [2025] show frontier reasoning models hacking verifiable rewards and obfuscating under monitoring; Krakovna et al. [2020] catalogue specification gaming — our entry is ordinary, the trace the novelty.

Shortcut learning, leakage, and validity. What the leaked policy learned is a shortcut [Geirhos et al., 2020]: a rule that succeeds in training because an input feature — word order — carried the label. The nearest analogues are NLI’s annotation artifacts — hypothesis-only models that pass [Gururangan et al., 2018, Poliak et al., 2018], heuristics that masquerade as inference [McCoy et al., 2019] — and Lapuschkin et al. [2019]’s Clever Hans; Kapoor and Narayanan [2023] taxonomize leakage for ML-based science. None usually reaches RL post-training, where the leak hides in code (§10) — our transfer. Measurement-theoretic critiques [Jacobs and Wallach, 2021, Freiesleben and Zezulka, 2025, Chehbouni et al., 2025, Manheim and Garrabrant, 2018] distinguish a measurement from its construct; ours is maximally reliable and the validity gap does the damage anyway.

12 Limitations

The leaked-SFT cell (§8) is one seed; every run shares one task and reward; no human baseline exists (only “worse than untrained” is established). An endpoint cannot rule out a traversed 1.5B peak (§4); the leaked peak may move across seeds (§3); a meaningful KL penalty is untested; the 1.5B loader shares the leak, copy rate unmeasured; the judges are correlated scorers (§9).

References

- Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021. Outstanding Paper. arXiv:2108.13264.
- Bowen Baker, Joost Huizinga, Leo Gao, Zehao Dou, Melody Y. Guan, Aleksander Madry, Wojciech Zaremba, Jakub Pachocki, and David Farhi. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*, 2025.
- Khaoula Chehbouni, Mohammed Haddou, Jackie Chi Kit Cheung, and Golnoosh Farnadi. Neither valid nor reliable? investigating the use of LLMs as judges. In *Advances in Neural Information Processing Systems 38 Main Conference (NeurIPS 2025)*, 2025. arXiv:2508.18076, 25 Aug 2025.
- Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, Zhiyuan Liu, Hao Peng, Lei Bai, Wanli Ouyang, Yu Cheng, Bowen Zhou, and Ning Ding. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025. Submitted 28 May 2025.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, et al. Deepseek-r1: Incentivizing reasoning capability in LLMs via reinforcement learning. *Nature*, 645:633–638, 2025. doi: 10.1038/s41586-025-09422-z. arXiv:2501.12948, 22 Jan 2025.
- Timo Freiesleben and Sebastian Zenz. The benchmarking epistemology: Construct validity for evaluating machine learning models. *arXiv preprint arXiv:2510.23191*, 2025. Submitted 27 Oct 2025.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 10835–10866. PMLR, 23–29 Jul 2023. arXiv:2210.10760, 19 Oct 2022.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.
- Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of NAACL-HLT*, 2018.
- Abigail Z. Jacobs and Hanna Wallach. Measurement and fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’21)*. Association for Computing Machinery, 2021. doi: 10.1145/3442188.3445901. arXiv:1912.05511.
- Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9):100804, 2023.
- Jacek Karwowski, Oliver Hayman, Xingjian Bai, Klaus Kiendlhofer, Charlie Griffin, and Joar Skalse. Goodhart’s law in reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2024.
- Guneet Kohli. Nine judges, two effective votes: Correlated errors undermine LLM evaluation panels. *arXiv preprint arXiv:2605.29800*, 2026. Submitted 28 May 2026.
- Victoria Krakovna, Jonathan Uesato, Vladimir Mikulik, Matthew Rahtz, Tom Everitt, Ramana Kumar, Zac Kenton, Jan Leike, and Shane Legg. Specification gaming: the flip side of AI ingenuity. DeepMind blog, 2020.
- Sebastian Lapuschkin, Stephan Wäldchen, Alexander Binder, Grégoire Montavon, Wojciech Samek, and Klaus-Robert Müller. Unmasking Clever Hans predictors and assessing what machines really learn. *Nature Communications*, 10(1):1096, 2019.

418 Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min
419 Lin. Understanding R1-Zero-like training: A critical perspective. *arXiv preprint arXiv:2503.20783*,
420 2025. Submitted 26 Mar 2025. Argues against group-level reward normalization.

421 David Manheim and Scott Garrabrant. Categorizing variants of Goodhart’s law. *arXiv preprint*
422 *arXiv:1803.04585*, 2018.

423 R. Thomas McCoy, Ellie Pavlick, and Tal Linzen. Right for the wrong reasons: Diagnosing syntactic
424 heuristics in natural language inference. In *Proceedings of ACL*, 2019.

425 Alexander Pan, Kush Bhatia, and Jacob Steinhardt. The effects of reward misspecification: Mapping
426 and mitigating misaligned models. In *International Conference on Learning Representations*
427 (*ICLR*), 2022.

428 Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme.
429 Hypothesis only baselines in natural language inference. In *Proceedings of the Seventh Joint*
430 *Conference on Lexical and Computational Semantics (*SEM)*, 2018.

431 Rulin Shao, Shuyue Stella Li, Rui Xin, Scott Geng, Yiping Wang, Sewoong Oh, Simon Shaolei Du,
432 Nathan Lambert, Sewon Min, Ranjay Krishna, Yulia Tsvetkov, Hannaneh Hajishirzi, Pang Wei
433 Koh, and Luke Zettlemoyer. Spurious rewards: Rethinking training signals in RLVR. *arXiv*
434 *preprint arXiv:2506.10947*, 2025.

435 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
436 Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of
437 mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. Introduces GRPO. Submitted 5 Feb 2024; v3 27 Apr 2024.

439 Joar Skalse, Nikolaus Howe, Dmitrii Krashennnikov, and David Krueger. Defining and characterizing
440 reward gaming. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022.
441 Published as “Reward Gaming”; circulated on arXiv as “Defining and Characterizing Reward
442 Hacking”, arXiv:2209.13085.

443 Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford,
444 Dario Amodei, and Paul F. Christiano. Learning to summarize with human feedback. In *Advances*
445 *in Neural Information Processing Systems 33 (NeurIPS 2020)*, 2020. Circulated on arXiv as
446 “Learning to summarize from human feedback”, arXiv:2009.01325.

447 Graham Todd, Tim Merino, Sam Earle, and Julian Togelius. Missed connections: Lateral thinking
448 puzzles for large language models. *arXiv preprint arXiv:2404.11730*, 2024.

449 Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang.
450 Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model?
451 In *Advances in Neural Information Processing Systems 38 Main Conference (NeurIPS 2025)*, 2025.
452 Oral. arXiv:2504.13837, 18 Apr 2025.

453 A The two prompts

454 The bug of §2 is one line, and it is easiest to see as the pair of strings the model actually received.
455 Both are the user turn of the same puzzle (id 3, 2023-06-14), under the same system prompt; the
456 answer key is {CHEEK, EYE, MOUTH, NOSE}, {CHOW, GOBBLE, SCARF, WOLF}, {LAB, PEKE,
457 PIT, POM}, {AMIGO, KING, STOOGE, TENOR}.

458 **GRPO training prompt, as the reported run consumed it** (train/grpo.py::build_dataset):

459 Words: CHEEK, EYE, MOUTH, NOSE, CHOW, GOBBLE, SCARF, WOLF,
460 LAB, PEKE, PIT, POM, AMIGO, KING, STOOGE, TENOR

461 **SFT data and every evaluation** (data/formatting.py::shuffled_words, a puzzle-seeded
462 permutation):

463 Words: CHOW, POM, PIT, KING, AMIGO, NOSE, SCARF, CHEEK,
 464 EYE, TENOR, STOOGE, GOBBLE, MOUTH, LAB, PEKE, WOLF

465 In the first, emitting the words four at a time in the order given scores the full 1.6. In the second it
 466 scores what a uniformly drawn valid partition scores. Nothing else differs — same system prompt,
 467 same reward, same parser — and the difference is a missing call to the shuffle the rest of the codebase
 468 already used. The regression test we have since added asserts the two paths agree per puzzle, which is
 469 the property that would have caught this and the weaker “is not in answer order” assertion would not.

470 B Per-checkpoint values for the entropy sweep

471 Table 5 tabulates all eleven rows of the sweep artifact: the trajectory plotted in Figure 1
 472 (including the step-400 checkpoint, indistinguishable from step 403 at this precision), the
 473 structural column discussed in Section 5, and the off-trajectory untrained Instruct reference.
 474 All values are from the $n = 100$ temperature-0.9 entropy sweep of the reported 7B run
 475 (seed 0); the artifact is `results-analysis/entropy-kl-7b.json` in the released repository,
 476 [repository URL withheld for review], which also holds the training and evaluation con-
 477 figs, the reward implementation, and the generation logs behind every table here. The valida-
 478 tion grid behind the step-50 selection (§3), artifact `results-analysis/ckpt-curve-7b.json`:
 479 greedy semantic score, 108 validation puzzles, steps 50/100/150/200/250/300/350/400/403 read
 480 **0.139/0.120/0.012/0.012/0.012/0.012/0.009/0.009/0.009** (SFT init 0.102); step 50 is the argmax of
 481 the nine, frozen before any test scoring.

482 Unless stated otherwise, all intervals in this paper are percentile bootstrap over *puzzles* — the unit
 483 of independent sampling — with $B = 1,000$ resamples at seed 0 (practice motivated by Agarwal
 484 et al., 2021); per-slot counts are reported alongside means but never resampled independently, since a
 485 puzzle’s four slots are not independent. Where a count is small enough that the bootstrap degenerates
 486 (Table 6), we substitute Clopper–Pearson intervals and say so.

Table 5: The reported 7B run on the $n = 100$ temperature-0.9 entropy sweep. Structural = fraction of outputs that are a valid partition; semantic = mean correct groups / 4. KL is the per-sequence sample-based estimate against the SFT initialization, each row estimated under its own policy’s samples; the untrained Instruct row is off the GRPO trajectory (§5).

Checkpoint	Policy entropy	KL from SFT init	Structural	Semantic
Instruct (untrained)	0.2577	42.60	0.69	0.0125
SFT init	0.3025	0.00	0.37	0.0325
Step 50	0.2107	2.70	0.68	0.0950
Step 100	0.1999	8.12	0.64	0.0875
Step 150	0.0158	46.42	0.95	0.0075
Step 200	0.0109	46.88	0.95	0.0100
Step 250	0.0102	46.96	0.96	0.0100
Step 300	0.0099	47.05	0.96	0.0100
Step 350	0.0099	47.05	0.96	0.0100
Step 400	0.0099	47.05	0.96	0.0100
Step 403	0.0099	47.05	0.96	0.0100

487 C 7B endpoints across three seeds

488 Table 6 tabulates the three-seed endpoint session discussed in Section 3; the 4 / 7 / 11 spread is the
 489 basis of the single-seed caveats in Sections 10 and 12. The pass@16 contrasts referenced there are
 490 differences of +0.920 [0.778, 1.062] for SFT over GRPO and −0.370 [−0.475, −0.259] for GRPO
 491 against the untrained Instruct model — single-seed, read as illustrative only.

Table 6: 7B endpoints across three seeds (separate serving session from Table 2; the SFT arm decodes to 0.321/52 both here and in the Table 2 control session, and to 0.346/56 in a third session; the untrained and final-checkpoint anchor arms reproduce across sessions exactly).

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward
Instruct (untrained)	0.160	26	6.8% [3.1, 11.1]	0.165
SFT	0.321	52	22.2%	0.197
GRPO (3 seeds)	0.045 ± 0.022	4 / 7 / 11	$1.2 \pm 0.6\%$	0.132 ± 0.008

D Control-session copy rates and rewards

Table 7 gives the full copy-rule test behind Section 7: for each generation, an emitted group counts as a copy when it equals, as a set, a consecutive quadruple (positions 1–4, 5–8, 9–12, 13–16) of that prompt’s word order. Train prompts parsed: 806 of 807 per arm; test: 162 of 162. Mean held-out rewards in this session: base 0.1802 (train) / 0.1654 (test), SFT 0.2490 / 0.1910, step 50 0.3113 / 0.2519, endpoint 0.1242 / 0.1250 — the endpoint’s test values reproduce the published 0.125 exactly. The per-stratum relative loss cited in §3 (untrained model → endpoint, over the 158 of 162 test puzzles with a single stratum label): 92% on category ($n = 69$), 67% on cultural ($n = 28$), 86% on tag-fillin ($n = 40$), 100% on wordplay ($n = 21$, from a near-floor base of 0.048, one recovered slot).

Table 7: Copy-rule rates, control session (greedy; shuffled evaluation prompts). Group-level = share of emitted groups that are prompt-order quadruples; response-level = share of responses whose four groups all are.

Arm	Train (group)	Train (response)	Test (group)	Test (response)
Instruct (untrained)	4.9% (159/3224)	4.3%	2.6%	1.2%
SFT	0.3% (9/3224)	0.0%	0.3%	0.0%
GRPO, step 50	0.3% (11/3224)	0.0%	0.2%	0.0%
GRPO, step 403	91.4% (2946/3224)	88.2%	92.4%	90.7%

E Leak-free rerun: checkpoint curve and training signals

Validation checkpoint curve of the §8 rerun ($n = 108$ validation puzzles, greedy; artifact results-analysis/aug29/ckpt-curve-7b-shuffled.json, evaluation outputs under results-analysis/aug29/leakfree-session-test/, and the full training log alongside them):

Step	Structural	Semantic (0–1)	Solve rate	Mean reward
50	0.852	0.1065	0.000	0.229
100	0.954	0.1574	0.028	0.331
150	0.972	0.1782	0.019	0.356
200	0.963	0.1736	0.009	0.351
250	0.954	0.1806	0.028	0.368
300	0.944	0.1944	0.028	0.376
350	0.963	0.1921	0.028	0.377
400	0.972	0.1921	0.028	0.381
403	0.972	0.1921	0.028	0.380

Training-signal summary over the 160 logged rows (single T4, QLoRA, 1.8×10^4 s for 403 steps): reward min/mean/max 0.123/0.371/0.705, rising $0.30 \rightarrow 0.44$ by halves; frac_reward_zero_std max 0.2, mean 0.031; across-rollout reward std is never zero on any logged row; trainer-logged policy entropy 0.116–0.212 and KL at most 0.22 (trainer instruments, not directly comparable to the §5 sweep estimates). One nonsubstantive difference from the leaked run: logging cadence (5-step averages there, 160 rows over 403 steps here), so cross-run in-sample extrema are extrema over differently smoothed series.

Seeds 1–2 (artifacts under results-analysis/sep05/): validation semantic score at steps 50/100/110/120/130/140/150/200/250/300/350/400/403 reads 0.141/0.160/0.162/0.157/0.171/0.190/0.194/**0.227**/0.215/0.220/0.218/0.220/0.220 for seed 1 and 0.130/0.162/0.169/0.174/0.185/0.185/0.197/0.199/**0.206**/0.183/0.185/0.190/0.194 for seed 2 — smooth climbs through the densely checkpointed 100–150 window, validation peaks at steps 200 and 250, plateaus thereafter. Test evaluations of those peaks read 127/648 and 132/648 against base’s 26 in their own session — close to their finals, though the two sit in different serving sessions, so no paired contrast is available. The leaked-SFT arm’s evaluation and the seed finals share the sep05 session with base and SFT anchors at 26 and 54 of 648.

F 1.5B test-split table

Table 8 tabulates the 1.5B arms discussed in Section 4; every mean is the exact ratio of its count to 162.

Table 8: 1.5B on the test split (162 puzzles, 648 group slots). Groups-correct intervals are Clopper–Pearson exact binomial on the slot count of 648, scaled to 0–4 (a percentile bootstrap is degenerate at counts of 1–2); slots cluster within puzzles, so they are approximate. Invalid-rate intervals are percentile bootstrap over puzzles. Conditioned on structurally valid output (§3), the arms read: base 0.009 ($n = 110$ of 162), SFT 0.048 ($n = 42$), GRPO 0.006 ($n = 158$).

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward	Solves
Base	0.0062 [0.0002, 0.0343]	1	32.1% [24.7, 38.9]	0.049	0.0%
SFT	0.0123 [0.0015, 0.0444]	2	74.1% [67.3, 80.2]	−0.038	0.0%
GRPO	0.0062 [0.0002, 0.0343]	1	2.5% [0.6, 5.6]	0.113	0.0%

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract’s numbers ($26 \rightarrow 52 \rightarrow 80 \rightarrow 4$, $+0.161/+0.173$, -0.136 , $91.4\%/92.4\%$, 21 of 3,228) are §3 and §7’s measured results; the diagnosis is stated as post hoc where it is one, and §12 bounds the scope claimed (three repaired seeds; the leaked-SFT cell is a single seed).

2. Limitations

Question: Have the authors discussed the limitations of their work?

Answer: [Yes]

Justification: §12 is a dedicated section and leads with the largest remaining limitations (single-seed cells; one task); it also covers the single-task design, the single-seed trajectory, the absent human baseline, correlated judges, and the untested KL regime.

3. Theory, Assumptions and Proofs

Question: Did you state the full set of assumptions of all theoretical results, and include complete proofs?

Answer: [N/A]

Justification: No theoretical results. The 1.4-slot chance floor is an exact combinatorial count ($4 \times 5775/2,627,625$), stated where used (§3).

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results?

Answer: [Yes]

Justification: The released repository ([repository URL withheld for review]) holds the training and evaluation configs, the reward implementation, the data-split construction, generation logs, and the control-session outputs behind every table; Appendix B names the artifact file for each sweep and the statistical conventions.

5. Open Access to Data and Code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results?

Answer: [Yes]

Justification: The repository ([repository URL withheld for review], withheld for double-blind review) contains code, configs, the evaluation harness with one-command entry points, and the regression test for the disclosed bug; adapters and result logs are released with documentation cards.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [Yes]

Justification: §2 gives the pipeline (rank-16 LoRA, three SFT epochs, $K = 8$ at temperature 0.9, $\beta = 0.001$, learning rate 5×10^{-6} , 403 optimizer steps); decoding and serving sessions are named per table by rule; the exact configs are in the repository.

7. Experiment Statistical Significance

Question: Does the paper report error bars or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All intervals are percentile bootstrap over puzzles (the unit of independent sampling), $B = 1,000$, seed 0, stated once in Appendix B; degenerate counts use Clopper–Pearson and say so; contrasts are paired within single serving sessions.

8. Experiments Compute Resources

Question: Does the paper provide sufficient information on the compute resources needed to reproduce the experiments?

Answer: [Yes]

Justification: All training ran in Kaggle notebooks on a single NVIDIA T4 (16 GB) via

582 QLoRA (4-bit base model, rank-16 LoRA adapters, gradient checkpointing); the 403-step
583 7B GRPO run takes about five hours (1.8×10^4 s measured for the leak-free rerun of the
584 identical recipe), and checkpoints sync to the Hub every 50 steps so a run approaching the
585 12-hour session cap resumes on a fresh session. The three seed runs and the normalization
586 ablation use the identical recipe and cost. Serving and evaluation ran on 2×T4 (16 GB) with
587 tensor parallelism under vLLM.

588 9. Code of Ethics

589 Question: Does the research conform, in every respect, with the NeurIPS Code of Ethics?

590 Answer: [Yes]

591 Justification: The work involves no human subjects, no personal data, and no high-risk
592 artifacts; the failure mode studied is disclosed in full, including the bug that caused it.

593 10. Broader Impacts

594 Question: Does the paper discuss both potential positive and negative societal impacts of
595 the work performed?

596 Answer: [N/A]

597 Justification: A measurement-validity case study on a word-puzzle task that releases no new
598 capabilities (the studied endpoint performs below the untrained base model on the construct).
599 Its practical import is an evaluation-practice recommendation (§10).

600 11. Safeguards

601 Question: Does the paper describe safeguards for responsible release of assets with high
602 risk for misuse?

603 Answer: [N/A]

604 Justification: The released models are LoRA adapters for NYT Connections; no high-misuse-
605 risk asset is released.

606 12. Licenses for Existing Assets

607 Question: Are the creators or original owners of assets used in the paper properly credited,
608 and are the license and terms of use explicitly mentioned and properly respected?

609 Answer: [Yes]

610 Justification: Qwen2.5 models (Apache 2.0) and the TRL/vLLM/PEFT stack (Apache 2.0)
611 are cited and used under their licenses. The words and answer keys of published NYT
612 Connections puzzles are held in a companion data repository whose data-provenance note
613 states their source (a public community archive of published puzzles), that The New York
614 Times retains all rights in the puzzles themselves, and that the repository’s MIT license is
615 scoped to its code only; the puzzle data is used for non-commercial research evaluation and
616 is not relicensed.

617 13. New Assets

618 Question: Are new assets introduced in the paper well documented, and is the documentation
619 provided alongside the assets?

620 Answer: [Yes]

621 Justification: The released code and configs are documented in the repository README;
622 released adapters and result sets carry documentation cards stating training provenance and
623 research-only intended use, including the disclosed prompt-order bug.

624 14. Crowdsourcing and Research with Human Subjects

625 Question: Does the paper include the full text of instructions given to participants, screen-
626 shots, and details about compensation?

627 Answer: [N/A]

628 Justification: No crowdsourcing and no human subjects.

629 15. IRB Approvals

630 Question: Does the paper describe potential risks to study participants and whether IRB
631 approvals were obtained?

632 Answer: [N/A]

633 Justification: No human subjects.

634 16. Declaration of LLM Usage

635 Question: Does the paper describe the usage of LLMs if it is an important, original, or
636 non-standard component of the core methods?

637 Answer: [Yes]

638 Justification: LLM judges are a declared component of the method: §9 names the three
639 models (gpt-5-mini, gemini-2.5-flash, claude-sonnet-4-5), the blind scoring protocol, and
640 their correlated-scorer limitation.