
Reliable and Invalid: Anatomy of a Gamed Training Task

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Verifiable rewards are trusted because they are deterministic: a program either
2 awards the points or it does not. We post-train Qwen2.5-7B-Instruct with GRPO on
3 NYT Connections under a reward with a cheap structural component (well-formed
4 output) and an expensive semantic component (correct groupings). On a strictly
5 chronological test set of 162 puzzles (648 group slots), scored in one serving
6 session, capability runs $26 \rightarrow 52 \rightarrow 80 \rightarrow 4$ slots across the untrained Instruct
7 model, supervised fine-tuning, GRPO step 50, and the GRPO endpoint: paired stage
8 gains of +0.161 and +0.173 on the 0–4 scale, then a fall to within $3\times$ of the 1.4-
9 slot chance floor of a uniformly drawn valid partition, $-0.136 [-0.198, -0.080]$
10 below the untrained model, reproduced across three seeds. The *training* reward
11 reads the destruction as mastery: it reaches its 1.6 ceiling by step 125 with zero
12 across-rollout variance and stays there. Scoring the endpoint offline on its own
13 training puzzles — 21 of 3,228 slots — rules out memorization and exposes the
14 cause: a one-line bug in our GRPO data loader presented each board’s sixteen
15 words in answer-key order, so the training task, as presented, was solvable by
16 copying the words four at a time. The collapsed policy does exactly that — 91.4%
17 of the groups it emits on training boards, and 92.4% on test boards, are consecutive
18 quadruples of its prompt; a copy of a shuffled board is a random valid partition,
19 which is why the endpoint lands near the chance floor. Entropy collapse, cross-seed
20 convergence, and the reward ceiling are signatures of one deterministic copy rule;
21 the structural climb is older, over half of it predating any measurable copying.
22 The reward function was correct on every string it scored; the task presentation
23 was gameable, and every in-sample signal read the gaming as mastery. The same
24 reward held out, and three blind LLM judges, catch the reversal; four lines of string
25 matching then identify the rule. Determinism certifies the scorer, not the construct:
26 hold the reward out.

27 1 Introduction

28 Post-training against verifiable rewards is attractive for a reason that sounds decisive: the reward is a
29 program. It returns exactly what it was written to return, and the argument that this buys immunity
30 from reward hacking¹ rests on that reliability. This paper is about the gap reliability does not close:
31 the gap between the reward and the construct it stands in for. We train on NYT Connections — sixteen
32 words, partitioned into four groups of four — chosen because well-formedness and correctness are
33 unusually separable: a response can be a perfectly valid partition and completely wrong. Our reward
34 pays for both; two of its five components can be checked by reading the string, three require the
35 answer key.

¹We follow the common usage *reward hacking* throughout; Skalse et al. [2022], whose formalization we adopt, publish it under the name *reward gaming*.

For fifty steps it works. In one serving session, held-out groups recovered run $26 \rightarrow 52 \rightarrow 80$ across the untrained Instruct model, the supervised checkpoint GRPO is initialized from, and step 50 — paired stage gains of $+0.161$ [$+0.056$, $+0.272$] and $+0.173$ [$+0.074$, $+0.284$], the reinforcement-learning stage contributing as much as supervised fine-tuning (§7). Over the following 353 steps the same run falls to 4 slots, below the untrained model and onto the 1.4-slot floor of a uniformly drawn valid partition, while the invalid-output rate falls by an order of magnitude and the training reward reaches its 1.6 maximum and averages 1.5814 thereafter, with zero variance across the $K = 8$ rollouts on $\sim 95\%$ of prompts. A practitioner watching that curve saw a solved task. The curve was not lying about the reward; it was lying about the task. A control run scoring the endpoint on its own training puzzles returned 21 of 3,228 slots — not memorization — and led to the mechanism: our GRPO data loader, unlike our SFT and evaluation code, presented each board’s words in answer-key order, and the converged policy simply copies the prompt four words at a time (91–92% of emitted groups; §7). The failure is not in the reward function, which is correct; it is in what the training distribution let the reward certify, and in the practice of reading the certificate in sample.

Contributions. A fully verifiable case study in which GRPO improves and then destroys the target construct, ending below the untrained model across three seeds (§3); a complete post-hoc diagnosis: an answer-ordered prompt bug made the training task gameable, and the converged policy is a prompt-order copier on 91–92% of its emitted groups (§7); an account in which entropy collapse, the reward ceiling, and cross-seed convergence are signatures of that one rule (§5), surviving a normalization ablation (§6); and the lesson demonstrated end to end: every in-sample signal read gaming as mastery; the reward held out, and three blind judges, caught it (§7–§8).

2 Task, pipeline, and the reward’s asymmetry

Task and data. 1,078 dated NYT Connections puzzles (1,077 after validation), split strictly chronologically: 807 train (2023-06-12 to 2025-08-27), 108 validation, 162 test (2025-12-15 to 2026-05-29); every test puzzle postdates every training puzzle. Date regressions of groups-correct across the test window are flat for both base (-0.158 groups/year [-0.634 , $+0.292$]) and the trained policy ($+0.035$ [-0.271 , $+0.329$]): contamination controlled, though the intervals are wide and only large drift would show.

Pipeline. Qwen2.5-7B-Instruct (and 1.5B for the scale contrast), rank-16 LoRA. SFT for 3 epochs on the training split, then GRPO [Shao et al., 2024] from the SFT checkpoint: $K = 8$ sampled completions per puzzle at temperature 0.9, group-relative advantages, $\beta = 0.001$ against the SFT reference, learning rate 5×10^{-6} , one epoch (403 optimizer steps) at 7B.

The leak, disclosed up front. SFT data and every evaluation in this paper present each board’s sixteen words deterministically shuffled (seeded by puzzle id), precisely so group order is not a giveaway. Our GRPO loader did not: a one-line bug built training prompts from the answer key, so during reinforcement learning — only there — the words arrived grouped, positions 1–4 the first group, and so on: every GRPO prompt was solvable by a positional copy worth the full 1.6. Discovered post hoc (§7); disclosed here because nothing in the training-time sections can be read without it. Held-out measurements are unaffected: evaluation prompts were always shuffled.

The reward. Table 1 is the crux of the paper.

Table 1: The reward decomposition. Two components are verifiable from the emitted string alone; three require the answer key. The policy is trained on the sum.

Component	Value	Cost to check
Parses as four groups of four	$+0.1$	cheap — read the string
Correct groups	$+1.0 \times (\text{correct}/4)$	needs the answer key
All four groups correct	$+0.5$	needs the answer key
Group with exactly 3 correct words	$+0.05$	needs the answer key
Invalid output	-0.1	cheap — read the string

Maximum 1.6. The model can raise its score by getting better at the puzzle or by getting better at emitting well-formed output; only one is the construct. That asymmetry — cheap and expensive components, independently satisfiable — is the design property this paper studies.

3 The arc on held-out data

Table 2 reports all four arms of the reported 7B run on the 162-puzzle test set, greedy decoding, served in a single vLLM session so every comparison is within-session. The supervised checkpoint is included because it, not the untrained model, is what GRPO was initialized from.

Table 2: 7B on the test split (162 puzzles, 648 group slots; all four arms in one serving session; greedy). Means on the 0–4 scale are exact ratios of the count column. Bootstrap intervals are percentile over puzzles, $B = 1,000$ resamples, seed 0. This is the session §7 also scores on the training split.

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward	Solves
Instruct (untrained)	0.1605 [0.105, 0.222]	26	6.8%	0.165	0/162
SFT	0.3210 [0.216, 0.432]	52	22.2%	0.191	2/162
GRPO, step 50	0.4938 [0.383, 0.623]	80	16.0%	0.252	2/162
GRPO, step 403	0.0247 [0.000, 0.056]	4	0.6%	0.125	0/162

Two reference points frame the table. A uniformly drawn valid partition recovers $4 \times 5775/2,627,625 = 0.0088$ groups per puzzle, or **1.4 slots** across this split; the step-403 endpoint sits at 2.8 times that floor — and §7 shows why: its copied responses score at chance, and the excess over the floor rides on the few responses that are not copies. Second, because all four arms sit in one session, the stage accounting is paired: $26 \rightarrow 52 \rightarrow 80 \rightarrow 4$, the supervised stage contributing $+0.1605 [+0.0556, +0.2716]$ and the reinforcement-learning stage $+0.1728 [+0.0741, +0.2840]$ — near-equal gains, both intervals excluding zero.

The method works, then reverses. Both movements are paired within Table 2’s session: the climb to step 50 is the two stage gains just given, each excluding zero, and step 403 falls below the untrained model by $-0.1358 [-0.1975, -0.0802]$. The earlier three-arm session read $+0.315 [+0.204, +0.438]$ and $-0.136 [-0.191, -0.080]$ for the same two contrasts; means in Table 2 are exact ratios of the count column. Two details locate the mechanism. Held-out mean reward reads 0.252 at step 50 and 0.125 at step 403 against base’s 0.165 (a separate session reads step 100 at 0.2664 to a step-50 re-serve’s 0.2525), while the training reward reaches its ceiling by step 125 (§5): the same reward ranks the checkpoints correctly out of sample and fails to on the training curve — that contrast, not the reward’s determinism, is this paper’s subject. And the step-50 policy is worse than base on structure (16.0% invalid against 6.8%) but *not* worse than its actual initialization, the SFT arm at 22.2% — so what is specific to the phase after step 50 is the semantic collapse, not the structural climb, which was already more than half complete before the peak (§5).

How wide is the peak? In a separate serving session, step 100 evaluates at 76/648 against a step-50 re-serve’s 80/648, paired difference $-0.0247 [-0.1481, +0.0927]$ ($n = 162$). We read this as uninformative rather than as equivalence: the interval’s width is comparable to the entire step-50 effect, so it is equally consistent with step 100 matching the peak and with step 100 having lost a third of it, and we did not pre-specify an equivalence margin. The normalization ablation of §6, this paper’s only independent trajectory, argues for the second reading — step 100 reads 17/400 against step 50’s 35/400 on the entropy sweep and 35/432 against 62/432 on the greedy checkpoint sweep, roughly half the peak gone in both. What both sessions agree on is narrower and is all we claim: the collapse falls between steps 100 and 150, where the validation semantic score drops $35/400 \rightarrow 3/400$ (§5).

Reproducibility across serving sessions. Base and the final checkpoint reproduce per-puzzle *exactly* across three vLLM sessions (0/162 differences each — a deterministic copier and a converged base leave no argmax ties for batched inference to resolve). Step 50 does not: 2/162 puzzles across sessions (77 vs. 80 of 648), 6/162 within one — mid-training policies sit near ties, which is why every step-50 number in this paper carries its serving session. Table 2’s 80/648 is the four-arm session used

117 throughout; an earlier three-arm session read 77/648, and the SFT arm decodes to 52 there and 56 in
118 a third (Appendix C). No conclusion in this paper turns on the difference.

119 **Robustness of the endpoint.** Three independent 7B seeds land at 4, 7 and 11 recovered slots of
120 648 (0.045 ± 0.022 on the 0–4 scale) — all below base’s 26. Their LoRA update directions agree
121 pairwise (cosine +0.67 to +0.69, merged weight deltas) — read descriptively, three runs finding the
122 same attractor, the copy rule (§7); seeds sharing an initialization and dataset are not isotropic draws,
123 so no random-direction null applies. Best-of-16 sampling does not rescue the converged policy: a
124 copier has nothing new to sample (single-seed, illustrative contrasts in Appendix C).

125 **Conditioning on well-formed output widens the gap rather than explaining it, and we label the**
126 **operation for what it is.** Is the effect an artifact of validity, since invalid generations score zero
127 groups and validity is what GRPO improved? Restricting to structurally valid generations at 7B, each
128 arm’s conditional computed within a single serving session (the SFT arm’s is its 56-slot session):
129 base **0.172** ($n = 151$ of 162), SFT **0.444** ($n = 126$), GRPO **0.025** ($n = 161$). This is a rescaling and
130 we label it as one: each conditional is the arm’s unconditional mean divided by its validity rate, so
131 the gap widens automatically when the low-scoring arm is the more valid, as here. It licenses only
132 the negative claim, the one we need: the deficit survives with formatting held fixed. Among answers
133 that parse, the converged policy recovers 4 group slots where the untrained model recovers 26.

134 **The loss shows no stratum gradient we can resolve.** Relative loss from the untrained model to
135 the endpoint runs 67–100% across all four stratum labels (values in Appendix D). We registered this,
136 pre-diagnosis, as a failed prediction of a selective-inference story; post-diagnosis it is the copy rule’s
137 expected signature, since a positional rule is blind to stratum (§7). The comparison cannot do more
138 than that: the strata hold 21–69 test puzzles each, the untrained model’s per-stratum base rates run as
139 low as one recovered slot, and we report no intervals — a design with no power to separate uniform
140 loss from a gradient. Read it as consistent with stratum-blindness, not as evidence for it.

141 4 The scale contrast at 1.5B

142 At 1.5B the instrument reads nothing. The untrained model recovers one group slot of 648, GRPO one,
143 and SFT — the best arm — two, against Section 3’s 1.4-slot chance floor. No arm is distinguishable
144 from that floor (SFT’s two slots carry an exact binomial interval of $[0.0015, 0.0444]$, Appendix E,
145 which contains it), so no claim about the construct is available at this scale. What is visible is
146 the reward driving its verifiable component: invalid output $32.1\% \rightarrow 2.5\%$, mean reward $0.049 \rightarrow$
147 0.113 (Appendix E). We report this as an observation: the runs are epoch-asymmetric (806 steps,
148 two epochs, against 403 and one), we have no 1.5B checkpoint trajectory, and “the construct was
149 untouched at 1.5B” would be an endpoint claim, exactly what this paper shows cannot settle the
150 question.

151 5 Trajectory of the collapse: entropy, KL, and the training signal

152 Figure 1 scores checkpoints of the reported run on the $n = 100$ temperature-0.9 *entropy sweep*
153 (denominators of 400; full per-checkpoint values in Appendix B); the ablation’s *checkpoint sweep*
154 (§6) is $n = 108$, greedy, /432 — different measurements, never mixed.

155 The semantic peak has a location; the structural climb does not. Semantic rises $0.0325 \rightarrow 0.0950$
156 from the SFT initialization to step 50 (13/400 $\rightarrow$ 38/400 slots), falls to 0.0075 by step 150, and holds
157 at 0.0100 from 200 on. Structural (valid-partition fraction, not shown) climbs $0.37 \rightarrow 0.68 \rightarrow 0.64$
158 $\rightarrow 0.95$ – 0.96 across init/50/100/150 on — one reversal, at 100 — with 0.31 of its 0.59 total gain
159 banked by step 50, concurrent with the semantic peak.

160 The trajectory is not a trade with a turning point: a broadly monotone structural climb (one reversal,
161 at 100) with a semantic hump on top, and only the hump has a location. One gap qualifies the shape:
162 no checkpoint exists between steps 100 and 150 — 81% of the KL displacement ($8.12 \rightarrow 46.42$ nats)
163 — so the descent is interpolated, not measured.

164 **Entropy collapse is the rule’s signature.** Policy entropy falls 13.4-fold between steps 50 and
165 150, the window in which the semantic score falls 12.7-fold; by step 150 the policy is effectively




Figure 1: The mechanism, on the $n = 100$ temperature-0.9 entropy sweep (reported 7B run, seed 0). **Left:** policy entropy against optimizer step — a 30-fold collapse, with the cliff between steps 100 and 150. **Right:** semantic score against KL from the SFT initialization: the peak arrives by 2.70 of an eventual 47.05 nats (5.7%; step 50 is the earliest checkpoint retained), the remainder undoes it. The untrained model is an off-trajectory reference at 42.60 nats: KL measures displacement, not damage (§5).

deterministic. With §7’s diagnosis in hand the reading is direct: a positional copy rule is deterministic, blind to the board’s content, and reliably well-formed: entropy collapse and semantic destruction read as two views of one convergence — a reading made after the diagnosis, not before it. The structural climb is not — over half of it ($0.37 \rightarrow 0.68$ by step 50) predates measurable copying (0.3% at step 50 by the strict all-four-quadruples rule of Appendix D; a graded partial-adherence statistic could date copying earlier, which is why we write *measurable*): format gains the rule later locks in. What the collapse removes is within-prompt variation — nothing new to sample (the pass@16 result), no group-relative learning signal (the zero-advantage fraction). And what survives on held-out boards is measured twice over: the cheap structural term ($0.1(1 - 2 \cdot \text{invalid})$, Table 2) is 52% / 27% / 79% of held-out reward across base / step 50 / endpoint, and the copy rule itself is visible in 90.7% of the endpoint’s test generations (92.4% of their emitted groups; §7).

KL is the budget view, and it is a path length, not a damage scale. The step-50 checkpoint sits 2.70 nats from the SFT initialization and is, jointly with step 100 (§3), the best policy this run produced on held-out data. The run spent 47.05: 80 slots at 5.7% of the displacement, 4 at all of it. Two qualifications keep this a path description, not a scale: step 50 is the earliest checkpoint retained, so “at 5.7%” means *at or before* it; and KL measures displacement, not damage — the untrained model sits 42.60 nats out on the same estimator (90.5% of the run’s displacement) while scoring 0.0125 on this sweep (Appendix B), each estimate under its own policy’s samples, so base is off this trajectory and 5.7% dates the peak within this run only. The sweep also compresses base’s margin over the endpoint (5/400 vs. 4/400 here; 26 vs. 4 on the greedy test instrument): temperature-0.9 sampling costs a high-entropy policy far more than a collapsed one.

The training signal, read from the trainer’s own logs. Figure 2 shows the paper’s central contrast on one axis: the training reward first touches the reward function’s ceiling of 1.6000 at step 125 and averages 1.5814 over the logged points from then on, with excursions that thin as training proceeds: the deepest is 1.4300 at step 155, the two halves of the 57 post-125 logged points average 1.5663 then 1.5960, and the count of points at exactly 1.6000 rises from 14 of 28 to 24 of 29. We describe this sharpening rather than test it — the series is autocorrelated and the metric bounded, so no trend statistic we computed is identifiable here, and part of the rise is the instrument (below). What the figure cannot show is TRL’s own `frac_reward_zero_std` diagnostic: the fraction of prompts whose eight sampled rewards are identical — zero group-relative advantage. It rises from 0.0 at step 50 to 1.0 at step 125 and averages 0.946 over the remainder of training, so from step 125 onward the optimizer receives no learning signal on nearly every prompt. The saturation is legible under the diagnosis: on answer-ordered prompts (§2) a copy rule earns 1.6 on every board, seen or unseen — consistent with saturation beginning mid-epoch, on first-visit batches, and with the vanishing variance. Read as diagnostics, both in-sample signals bracket the collapse — the mean’s climb to ceiling ($0.9931 \rightarrow 1.6000$) and the variance’s fall to zero both land between logged steps 100 and




Figure 2: The same reward, in sample and held out (reported 7B run, seed 0). Training reward (gray; sampled $T=0.9$, $K=8$; TRL 5-step averages; measured on the answer-ordered prompts of §2) first touches its 1.6 ceiling at step 125, averaging 1.5814 thereafter. Held-out mean reward (162 test puzzles; solid markers; greedy; the control session of Table 3; right axis, note the scale) peaks at 0.252 at step 50 and ends at 0.125, below the untrained model’s 0.165; the hollow step-100 point is a later session (0.2664, vs. a step-50 re-serve at 0.2525) and is not joined.

125 — and neither signs it: saturation with zero spread is what mastery and degeneracy alike look like. Part of the climb is the instrument: as entropy collapses, temperature-0.9 sampling converges on greedy, so some of the rise is the sampling penalty vanishing (the sweep-vs-test gap, reversed). The remaining 278 steps of movement while $\sim 95\%$ of prompts carry zero advantage are consistent with continued sharpening of the rule on the residual prompts that still do; the weight-space counterpart of that push we have not traced.

β was numerically inert, and the bracket is not an experiment. An earlier configuration at learning rate 10^{-6} and $\beta = 0.04$ froze the policy (KL ≈ 0.001 for 403 steps; greedy outputs byte-identical to SFT); the reported one, at 5×10^{-6} and $\beta = 0.001$, overshoots by roughly as much. Two qualifications, both weakening the pair as an experiment: the configurations differ in *two* variables at once (β by $40\times$, learning rate by $5\times$), so the pair bounds a region jointly and attributes nothing to either variable; and $\beta = 0.001$ was numerically inert — $\beta \cdot \text{KL} = 0.0027$ reward units at step 50, 0.047 at convergence — so the KL reference never meaningfully restrained the policy, this paper has no evidence about a KL penalty large enough to matter, and the intermediate- β sweep remains unrun.

6 The ablation: the collapse is not the normalization default

We inherited `scale_rewards="group"`, verified as the trl 1.10.0 default. Reviewers of an earlier draft independently named it as a confound: near determinism, the within-group standard deviation dividing the advantage approaches zero, amplifying sampling noise into large updates — which alone might produce the collapse, its timing, and the cross-seed agreement; Liu et al. [2025] argue against the normalization on independent grounds.

We reran the 7B pipeline identically — same SFT warm start, learning rate, β , K , temperature, LoRA rank, epoch count, seed, and the same answer-ordered prompts of §2, so this ablation isolates the normalization, not the leak — with reward scaling disabled, the resolved value read back off the constructed trainer configuration. The signature is unchanged on the greedy $n = 108$ checkpoint sweep: validation semantic score rises to 62/432 at step 50, falls to 35/432 by step 100 and 4/432 by step 150, while structural validity climbs 0.70 to 0.96. The validation-selected peak (step 50, frozen before any test puzzle was scored) recovers 87/648 on test, $+0.377$ [$+0.247$, $+0.506$] paired against base; the final policy 7/648, below base by -0.117 [-0.179 , -0.056] paired and within $+0.019$ of the original endpoint.

The signature survives in the dynamics too, though not point for point: without reward scaling, policy entropy falls 0.2111 (step 50) $\rightarrow$ 0.1346 (100) $\rightarrow$ 0.0128 (150) and stays near 0.011 through step 403

Table 3: The control session: four arms, one vLLM session, greedy, both splits, shuffled evaluation prompts; percentile bootstrap, seed 0. Copy rate = emitted groups equal to consecutive quadruples of the prompt’s word order (Appendix D).

Arm	Train groups (0–4)	of 3,228	Test groups (0–4)	of 648	Copy rate (test)
Instruct (untrained)	0.1871 [0.156, 0.224]	151	0.1605 [0.105, 0.222]	26	2.6%
SFT	0.5118 [0.461, 0.568]	413	0.3210 [0.216, 0.432]	52	0.3%
GRPO, step 50	0.6803 [0.616, 0.746]	549	0.4938 [0.383, 0.623]	80	0.2%
GRPO, step 403	0.0260 [0.012, 0.041]	21	0.0247 [0.000, 0.056]	4	92.4%

(original endpoint 0.0099), while KL rises $2.75 \rightarrow 17.84 \rightarrow 45.64$ nats per sequence and plateaus near 46 (original 47.05; 17.84 at step 100 vs. the original’s 8.12 — intermediate magnitudes differ). On the $n = 100$ temperature-0.9 entropy sweep — a different measurement from the greedy $n = 108$ checkpoint sweep above — semantic score reads 35/400 at step 50, 17/400 at 100, then 2/400 from step 150 onward. Hump, collapse, entropy endpoint, below-base endpoint: all survive with the normalization off. The phenomenon is not a library default.

7 The control run: not memorization — a copy rule

Training-reward saturation says the policy solves its training prompts under sampling; was that SFT-stage memorization? The control: all four arms served in one vLLM session, greedy, scored on the full 807-puzzle training split and the 162-puzzle test split under the same shuffled prompts as every other evaluation (Table 3).

Three results. The endpoint retains nothing — 21 of 3,228 training slots, indistinguishable from its own test floor: not a train/test gap, so not memorization. SFT does partially memorize, though less than the raw numbers suggest: its +0.191 train-minus-test gap (0.5118 vs. 0.3210 after three epochs) must be read against the untrained model’s own +0.027 gap on the same chronological splits (0.1871 vs. 0.1605), which is split difficulty rather than memorization; the adjusted excess is +0.164. Step 50 inherits essentially all of it (+0.160 adjusted) while improving on SFT by the same margin on both splits (+0.169 / +0.173): the peak’s gain generalizes. And the shared session yields the paired stage split: 26 $\rightarrow$ 52 $\rightarrow$ 80 $\rightarrow$ 4, SFT +0.1605 [+0.0556, +0.2716], RL +0.1728 [+0.0741, +0.2840], endpoint -0.1358 [-0.1975 , -0.0802] vs. the untrained model (the original session read -0.136 [-0.191 , -0.080]).

Four lines of string matching find the mechanism. A 1.6 training reward against 21 recovered training slots forced a look at the training prompts, where the bug of §2 was found. The behavioral test counts emitted groups equal, as sets, to consecutive quadruples of the prompt’s word order (definition and full counts in Appendix D). Such quadruples are 91.4% of the groups the endpoint emits on training boards and 92.4% on test boards; 88.2% and 90.7% of its responses are pure four-quadruple copies; SFT and step 50 sit at 0.3%, below even the untrained model’s weak bias (4.9% / 2.6%). The converged policy is a prompt-order copier. On training prompts, where order was the answer, copying earns 1.6; on shuffled prompts a copy is a fixed valid partition of a random order — exactly the uniformly drawn valid partition behind §3’s 1.4-slot floor. Its pure-copy responses score no better than that process: 2 training slots against 6.25 chance-expected — a three-fold deficit, if anything, and certainly not an excess (we attach no test to it: a puzzle’s four slots are not independent, Appendix B). What sits above chance is the aggregate — 21 slots against 7.1 expected; Table 3’s puzzle-level interval, [0.012, 0.041] on the 0–4 scale, excludes the chance mean of 0.0088, and test reads 4 against 1.4, same direction, interval not resolving — and it is carried almost entirely by responses that are *not* pure copies (35 non-copies recover 16 of the 21 slots; 60 partial copies, 3; 711 pure copies, 2); residual solving the rule has not displaced, not copying that works.

8 The reversal is legible from outside the reward

The 162 test generations from base, step 50 and the final policy were scored blind (board words and completion only) by three LLM judges (gpt-5-mini, gemini-2.5-flash, claude-sonnet-4-5) on a 1–10 scale, 1458/1458 responses parsed. All three rank step 50 above base above the final policy, all nine

pairwise per-puzzle differences excluding zero under a paired bootstrap (seed 0; step-50 minus final: $+0.60$ [$+0.44, +0.78$] / $+3.25$ [$+2.81, +3.67$] / $+1.62$ [$+1.33, +1.90$]). That ordering rules out the obvious confound: the final policy is the *most* well-formed arm on this split (0.6% invalid against 6.8% and 16.0%), so a shared preference for well-formed text would have placed it first, not last. In hindsight the judges were reading the copy rule: the arm they unanimously rank last is the copier (§7). Judges do err in correlated ways [Kohli, 2026], so we read this as three correlated readings of one external signal rather than three independent confirmations; per-item Pearson r is 0.456 / 0.587 / 0.626 across the judge pairs (486 items), lowest on the final arm (0.200–0.433); the nine comparisons are one uncorrected family, all excluding zero by margins no standard correction would close. One further limit: the judges score whether the groupings are right — the reward’s semantic component through a different instrument — so this corroborates the held-out reading rather than adding a second construct.

9 Implications for evaluation practice

Distinguish in-sample from held-out proxy metrics before concluding a proxy has failed. Evaluated on held-out data, our reward ranks the three checkpoints correctly; it fails only in the form practitioners actually consult it — as a training curve, in sample, where it reads perfect. The cheap remedy is not a better reward but the same reward on data the policy was not updated against — plus the trainer’s own in-sample brackets (mean saturation, zero variance), though neither signs the turn (§5). Held out, the separation is real but not large: paired against the untrained model over the 162 test puzzles, mean reward moves $+0.086$ [$+0.045, +0.129$] at step 50 and -0.040 [$-0.061, -0.019$] at the endpoint, both excluding zero (SFT, at $+0.026$ [$-0.015, +0.068$], does not). We flag the margin because 79% of the endpoint’s held-out reward is the cheap structural term: a reward weighted slightly differently could have kept the aggregate above base while the construct collapsed exactly as it did. The aggregate caught this run; the decomposition below is what makes it a diagnostic rather than a coincidence.

Determinism certifies the scorer, not the task. Our reward was correct on every string it scored, and the run was still gamed — through the prompt distribution. The audit cost four lines of string matching, but nothing in-sample prompted it: held-out evaluation said there was something to find. One function, two measurements: in sample, reliable and invalid — this paper’s title; held out, what validity it had.

Report reward components separately, never the aggregate. This is the recommendation we would keep if we could keep only one. The failure is invisible in the sum and obvious in the decomposition, which costs nothing: §5 reports it for our three arms (52%, 27%, 79%). It is a standard reward-model diagnostic, simply not applied to deterministic rewards: determinism is taken to have settled the question.

Treat single-seed RL results as uninterpretable. The caution applies to results that come out well too, including step 50: one seed (endpoint spread across seeds: 0.025–0.068), the argmax of a nine-point validation grid (step 100 reads 0.0875 to its 0.0950 at $n = 100$, inside one standard error). The grid preceded every test evaluation of the selected checkpoint — ordinary validation-based selection. Whether step 50 is a point or a shoulder we cannot say: §3’s step-100 evaluation is too wide to separate the two, and the ablation trajectory suggests decline has already begun by then.

10 Related work

Reward-model overoptimization. Gao et al. [2023] established this curve shape for *learned* reward models — gold-standard performance rising then declining with KL — where the explanation is Goodhart against an approximation [Stienon et al., 2020]. Our reward is a program: it approximates nothing, and there was nothing in the *function* to game — the gaming ran through the prompt distribution (§2), a channel the learned-reward literature’s remedy, a better reward model, does not touch. Determinism is the very feature cited to argue verifiable rewards escape this failure mode; here it guaranteed only that the shortcut was scored accurately.

RL for reasoning and entropy collapse. Verifiable-reward RL underlies recent reasoning systems [DeepSeek-AI et al., 2025, Shao et al., 2024]; Yue et al. [2025] ask whether it extends capability beyond the base model’s pass@ k ceiling, and here RLVR, initialized from an instruction-tuned stack, ends below *both* of its own antecedents on the construct. Shao et al. [2025] show Qwen2.5 models gaining on MATH-500 under RLVR with random or incorrect rewards, by amplifying latent base-model behaviors; part of our step-50 gain, obtained under a corrupted signal, may be that phenomenon — the peak is not evidence the signal was sound. Cui et al. [2025] document entropy collapse as a central dynamic of RL for reasoning; here the collapse coincides with convergence onto a leaked shortcut, and we read it as symptom rather than independent cause — a post-hoc reading the shuffled-prompt rerun can test (§7, §11). Liu et al. [2025] critique group-level advantage normalization; §6 shows the phenomenon does not depend on it.

Reward hacking and its taxonomy. Skalse et al. [2022] formalize the failure mode (as *reward gaming*); ours is gaming in their sense against the composed task — reward plus prompt distribution — not the reward function, which was exact. Abahana [2026] classifies RLHF failure transitions by directional changes in learned reward and judge scores; ours differs in instrument (a rule-based reward, exact) and in detection (bracketed by held-out data rather than flagged in flight). Krakovna et al. [2020] catalogue specification gaming across RL systems; the entry our run would occupy is an ordinary one, and we claim novelty for the trace rather than the phenomenon.

Shortcut learning and leakage. What our policy learned is a shortcut in the sense of Geirhos et al. [2020]: a decision rule that succeeds on the training distribution and fails off it, because a feature of the input — here, word order — carried the label. Lapuschkin et al. [2019] make the same point for supervised vision models under the name Clever Hans prediction. Neither literature is usually applied to RL post-training, where the training distribution is regenerated by a trainer, so a leak hides in code rather than in data. That is the transfer we are making.

Abruptness. Pan et al. [2022] report phase transitions in reward misspecification — capability thresholds at which true reward falls sharply — in some environments they study. Our collapse between steps 100 and 150 has that shape, which is why the gap in our checkpoint grid is a real limitation (§11), not an interpolation detail: if the turn is a transition, its location is the measurement worth having.

Validity in AI evaluation. Measurement-theoretic critiques of ML evaluation [Jacobs and Wallach, 2021, Freiesleben and Zezulka, 2025] distinguish a measurement from the construct it operationalizes; Chehbouni et al. [2025] apply it to LLM judges; ours is a case where the measurement is maximally reliable and the validity gap does the damage anyway. Agarwal et al. [2021] motivate our statistical practice.

11 Limitations

The largest limitation: no leak-free run exists. Every training run here (the reported 7B, the three seeds, the ablation) consumed the answer-ordered prompts, so nothing separates what GRPO does to this construct from what it did to the leaked presentation of it; the one-line fix is in the repository, and the shuffled-prompt rerun is the experiment these conclusions wait on. One task, one reward design; the checkpoint trajectory is 7B seed 0 only, so the scale claim rests on endpoints at both scales and a trajectory at one, and an endpoint cannot rule out a traversed-and-returned 1.5B peak (§4). No human baseline exists: “worse than the untrained model” is established, “bad in absolute terms” is not. Three-seed evidence covers only the converged endpoint (spread 0.025–0.068, why we distrust single numbers); the peak’s height and location may move across seeds, and our step-100 evaluation is too wide to establish that it is not a knife edge (§3). A single β is reported alongside one that froze the policy (no intermediate run, a learning-rate difference besides, $\beta \cdot \text{KL} \leq 0.047$ throughout), so a working KL penalty remains untested. The diagnosis is post hoc (§5’s readings are retrospective), the 1.5B loader shares the leak (copy rate unmeasured), and the blind judges are correlated scorers, not independent votes (§8).

References

- Zelalem Abahana. When RLHF fails: A mechanistic taxonomy of reward hacking, collapse, and evaluator gaming. *arXiv preprint arXiv:2606.03238*, 2026. Submitted 2 Jun 2026. Single-author preprint; no peer-reviewed version as of 24 Aug 2026.
- Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021. Outstanding Paper. arXiv:2108.13264.
- Khaoula Chehbouni, Mohammed Haddou, Jackie Chi Kit Cheung, and Golnoosh Farnadi. Neither valid nor reliable? investigating the use of LLMs as judges. In *Advances in Neural Information Processing Systems 38 Main Conference (NeurIPS 2025)*, 2025. arXiv:2508.18076, 25 Aug 2025.
- Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, Zhiyuan Liu, Hao Peng, Lei Bai, Wanli Ouyang, Yu Cheng, Bowen Zhou, and Ning Ding. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025. Submitted 28 May 2025.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, et al. Deepseek-r1: Incentivizing reasoning capability in LLMs via reinforcement learning. *Nature*, 645:633–638, 2025. doi: 10.1038/s41586-025-09422-z. arXiv:2501.12948, 22 Jan 2025.
- Timo Freiesleben and Sebastian Zenz. The benchmarking epistemology: Construct validity for evaluating machine learning models. *arXiv preprint arXiv:2510.23191*, 2025. Submitted 27 Oct 2025.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 10835–10866. PMLR, 23–29 Jul 2023. arXiv:2210.10760, 19 Oct 2022.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.
- Abigail Z. Jacobs and Hanna Wallach. Measurement and fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’21)*. Association for Computing Machinery, 2021. doi: 10.1145/3442188.3445901. arXiv:1912.05511.
- Guneet Kohli. Nine judges, two effective votes: Correlated errors undermine LLM evaluation panels. *arXiv preprint arXiv:2605.29800*, 2026. Submitted 28 May 2026.
- Victoria Krakovna, Jonathan Uesato, Vladimir Mikulik, Matthew Rahtz, Tom Everitt, Ramana Kumar, Zac Kenton, Jan Leike, and Shane Legg. Specification gaming: the flip side of AI ingenuity. DeepMind blog, 2020.
- Sebastian Lapuschkin, Stephan Wäldchen, Alexander Binder, Grégoire Montavon, Wojciech Samek, and Klaus-Robert Müller. Unmasking Clever Hans predictors and assessing what machines really learn. *Nature Communications*, 10(1):1096, 2019.
- Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min Lin. Understanding R1-Zero-like training: A critical perspective. *arXiv preprint arXiv:2503.20783*, 2025. Submitted 26 Mar 2025. Argues against group-level reward normalization.
- Alexander Pan, Kush Bhatia, and Jacob Steinhardt. The effects of reward misspecification: Mapping and mitigating misaligned models. In *International Conference on Learning Representations (ICLR)*, 2022.
- Rulin Shao, Shuyue Stella Li, Rui Xin, Scott Geng, Yiping Wang, Sewoong Oh, Simon Shaolei Du, Nathan Lambert, Sewon Min, Ranjay Krishna, Yulia Tsvetkov, Hannaneh Hajishirzi, Pang Wei Koh, and Luke Zettlemoyer. Spurious rewards: Rethinking training signals in RLVR. *arXiv preprint arXiv:2506.10947*, 2025.

422 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
423 Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of
424 mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. Intro-
425 duces GRPO. Submitted 5 Feb 2024; v3 27 Apr 2024.

426 Joar Skalse, Nikolaus Howe, Dmitrii Krashenninikov, and David Krueger. Defining and characterizing
427 reward gaming. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022.
428 Published as “Reward Gaming”; circulated on arXiv as “Defining and Characterizing Reward
429 Hacking”, arXiv:2209.13085.

430 Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford,
431 Dario Amodei, and Paul F. Christiano. Learning to summarize with human feedback. In *Advances
432 in Neural Information Processing Systems 33 (NeurIPS 2020)*, 2020. Circulated on arXiv as
433 “Learning to summarize from human feedback”, arXiv:2009.01325.

434 Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang.
435 Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model?
436 In *Advances in Neural Information Processing Systems 38 Main Conference (NeurIPS 2025)*, 2025.
437 Oral. arXiv:2504.13837, 18 Apr 2025.

438 A The two prompts

439 The bug of §2 is one line, and it is easiest to see as the pair of strings the model actually received.
440 Both are the user turn of the same puzzle (id 3, 2023-06-14), under the same system prompt; the
441 answer key is {CHEEK, EYE, MOUTH, NOSE}, {CHOW, GOBBLE, SCARF, WOLF}, {LAB, PEKE,
442 PIT, POM}, {AMIGO, KING, STOOGE, TENOR}.

443 **GRPO training prompt, as the reported run consumed it** (train/grpo.py::build_dataset):

444 Words: CHEEK, EYE, MOUTH, NOSE, CHOW, GOBBLE, SCARF, WOLF,
445 LAB, PEKE, PIT, POM, AMIGO, KING, STOOGE, TENOR

446 **SFT data and every evaluation** (data/formatting.py::shuffled_words, a puzzle-seeded
447 permutation):

448 Words: CHOW, POM, PIT, KING, AMIGO, NOSE, SCARF, CHEEK,
449 EYE, TENOR, STOOGE, GOBBLE, MOUTH, LAB, PEKE, WOLF

450 In the first, emitting the words four at a time in the order given scores the full 1.6. In the second it
451 scores what a uniformly drawn valid partition scores. Nothing else differs — same system prompt,
452 same reward, same parser — and the difference is a missing call to the shuffle the rest of the codebase
453 already used. The regression test we have since added asserts the two paths agree per puzzle, which is
454 the property that would have caught this and the weaker “is not in answer order” assertion would not.

455 B Per-checkpoint values for the entropy sweep

456 Table 4 tabulates all eleven rows of the sweep artifact: the trajectory plotted in Figure 1 (in-
457 cluding the step-400 checkpoint, indistinguishable from step 403 at this precision), the structural
458 column discussed in Section 5, and the off-trajectory untrained Instruct reference. All values
459 are from the $n = 100$ temperature-0.9 entropy sweep of the reported 7B run (seed 0); the arti-
460 fact is results-analysis/entropy-kl-7b.json in the released repository, [repository URL
461 withheld for review], which also holds the training and evaluation configs, the reward imple-
462 mentation, and the generation logs behind every table here.

463 Unless stated otherwise, all intervals in this paper are percentile bootstrap over *puzzles* — the unit of
464 independent sampling — with $B = 1,000$ resamples at seed 0; per-slot counts are reported alongside
465 means but never resampled independently, since a puzzle’s four slots are not independent. Where
466 a count is small enough that the bootstrap degenerates (Table 5), we substitute Clopper–Pearson
467 intervals and say so.

Table 4: The reported 7B run on the $n = 100$ temperature-0.9 entropy sweep. Structural = fraction of outputs that are a valid partition; semantic = mean correct groups / 4. KL is the per-sequence sample-based estimate against the SFT initialization, each row estimated under its own policy’s samples; the untrained Instruct row is off the GRPO trajectory (§5).

Checkpoint	Policy entropy	KL from SFT init	Structural	Semantic
Instruct (untrained)	0.2577	42.60	0.69	0.0125
SFT init	0.3025	0.00	0.37	0.0325
Step 50	0.2107	2.70	0.68	0.0950
Step 100	0.1999	8.12	0.64	0.0875
Step 150	0.0158	46.42	0.95	0.0075
Step 200	0.0109	46.88	0.95	0.0100
Step 250	0.0102	46.96	0.96	0.0100
Step 300	0.0099	47.05	0.96	0.0100
Step 350	0.0099	47.05	0.96	0.0100
Step 400	0.0099	47.05	0.96	0.0100
Step 403	0.0099	47.05	0.96	0.0100

C 7B endpoints across three seeds

Table 5 tabulates the three-seed endpoint session discussed in Section 3; the 4 / 7 / 11 spread is the basis of the single-seed caveats in Sections 9 and 11. The pass@16 contrasts referenced there are differences of +0.920 [0.778, 1.062] for SFT over GRPO and −0.370 [−0.475, −0.259] for GRPO against the untrained Instruct model — single-seed, read as illustrative only.

Table 5: 7B endpoints across three seeds (separate serving session from Table 2; the SFT arm decodes to 0.346/56 in the Table 2 session’s configuration and 0.321/52 here, differing on 2 of 162 puzzles; the untrained and final-checkpoint anchor arms reproduce across sessions exactly).

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward
Instruct (untrained)	0.160	26	6.8% [3.1, 11.1]	0.165
SFT	0.321	52	22.2%	0.197
GRPO (3 seeds)	0.045 ± 0.022	4 / 7 / 11	1.2 ± 0.6%	0.132 ± 0.008

D Control-session copy rates and rewards

Table 6 gives the full copy-rule test behind Section 7: for each generation, an emitted group counts as a copy when it equals, as a set, a consecutive quadruple (positions 1–4, 5–8, 9–12, 13–16) of that prompt’s word order. Train prompts parsed: 806 of 807 per arm; test: 162 of 162. Mean held-out rewards in this session: base 0.1802 (train) / 0.1654 (test), SFT 0.2490 / 0.1910, step 50 0.3113 / 0.2519, endpoint 0.1242 / 0.1250 — the endpoint’s test values reproduce the published 0.125 exactly. The per-stratum relative loss cited in §3 (untrained model → endpoint, over the 158 of 162 test puzzles with a single stratum label): 92% on category ($n = 69$), 67% on cultural ($n = 28$), 86% on tag-fillin ($n = 40$), 100% on wordplay ($n = 21$, from a near-floor base of 0.048, one recovered slot).

Table 6: Copy-rule rates, control session (greedy; shuffled evaluation prompts). Group-level = share of emitted groups that are prompt-order quadruples; response-level = share of responses whose four groups all are.

Arm	Train (group)	Train (response)	Test (group)	Test (response)
Instruct (untrained)	4.9% (159/3224)	4.3%	2.6%	1.2%
SFT	0.3% (9/3224)	0.0%	0.3%	0.0%
GRPO, step 50	0.3% (11/3224)	0.0%	0.2%	0.0%
GRPO, step 403	91.4% (2946/3224)	88.2%	92.4%	90.7%

483 **E 1.5B test-split table**

484 Table 7 tabulates the 1.5B arms discussed in Section 4; every mean is the exact ratio of its count to
 485 162.

Table 7: 1.5B on the test split (162 puzzles, 648 group slots). Groups-correct intervals are Clopper–Pearson exact binomial on the slot count of 648, scaled to 0–4 (a percentile bootstrap is degenerate at counts of 1–2); slots cluster within puzzles, so they are approximate. Invalid-rate intervals are percentile bootstrap over puzzles. Conditioned on structurally valid output (§3), the arms read: base 0.009 ($n = 110$ of 162), SFT 0.048 ($n = 42$), GRPO 0.006 ($n = 158$).

Arm	Groups correct (0–4)	Groups (of 648)	Invalid	Mean reward	Solves
Base	0.0062 [0.0002, 0.0343]	1	32.1% [24.7, 38.9]	0.049	0.0%
SFT	0.0123 [0.0015, 0.0444]	2	74.1% [67.3, 80.2]	−0.038	0.0%
GRPO	0.0062 [0.0002, 0.0343]	1	2.5% [0.6, 5.6]	0.113	0.0%

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract’s numbers ($26 \rightarrow 52 \rightarrow 80 \rightarrow 4$, $+0.161/+0.173$, -0.136 , $91.4\%/92.4\%$, 21 of 3,228) are §3 and §7’s measured results; the diagnosis is stated as post hoc where it is one, and §11 bounds the scope claimed (no leak-free run exists).

2. Limitations

Question: Have the authors discussed the limitations of their work?

Answer: [Yes]

Justification: §11 is a dedicated section and leads with the largest limitation (no leak-free run); it also covers the single-task design, the single-seed trajectory, the absent human baseline, correlated judges, and the untested KL regime.

3. Theory, Assumptions and Proofs

Question: Did you state the full set of assumptions of all theoretical results, and include complete proofs?

Answer: [N/A]

Justification: No theoretical results. The 1.4-slot chance floor is an exact combinatorial count ($4 \times 5775/2,627,625$), stated where used (§3).

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results?

Answer: [Yes]

Justification: The released repository ([repository URL withheld for review]) holds the training and evaluation configs, the reward implementation, the data-split construction, generation logs, and the control-session outputs behind every table; Appendix B names the artifact file for each sweep and the statistical conventions.

5. Open Access to Data and Code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results?

Answer: [Yes]

Justification: The repository ([repository URL withheld for review], withheld for double-blind review) contains code, configs, the evaluation harness with one-command entry points, and the regression test for the disclosed bug; adapters and result logs are released with documentation cards.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [Yes]

Justification: §2 gives the pipeline (rank-16 LoRA, three SFT epochs, $K = 8$ at temperature 0.9, $\beta = 0.001$, learning rate 5×10^{-6} , 403 optimizer steps); decoding and serving sessions are named per table by rule; the exact configs are in the repository.

7. Experiment Statistical Significance

Question: Does the paper report error bars or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All intervals are percentile bootstrap over puzzles (the unit of independent sampling), $B = 1,000$, seed 0, stated once in Appendix B; degenerate counts use Clopper–Pearson and say so; contrasts are paired within single serving sessions.

8. Experiments Compute Resources

Question: Does the paper provide sufficient information on the compute resources needed to reproduce the experiments?

Answer: [TODO]

Justification: [Author to confirm before submission: hardware and wall-clock for the

reported 7B GRPO run, the three seed runs, and the ablation. The control-session and copy-rule evaluations ran on 2×T4 16GB.]

9. Code of Ethics

Question: Does the research conform, in every respect, with the NeurIPS Code of Ethics?

Answer: [Yes]

Justification: The work involves no human subjects, no personal data, and no high-risk artifacts; the failure mode studied is disclosed in full, including the bug that caused it.

10. Broader Impacts

Question: Does the paper discuss both potential positive and negative societal impacts of the work performed?

Answer: [N/A]

Justification: A measurement-validity case study on a word-puzzle task that releases no new capabilities (the studied endpoint performs below the untrained base model on the construct). Its practical import is an evaluation-practice recommendation (§9).

11. Safeguards

Question: Does the paper describe safeguards for responsible release of assets with high risk for misuse?

Answer: [N/A]

Justification: The released models are LoRA adapters for NYT Connections; no high-misuse-risk asset is released.

12. Licenses for Existing Assets

Question: Are the creators or original owners of assets used in the paper properly credited, and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Qwen2.5 models (Apache 2.0) and the TRL/vLLM/PEFT stack (Apache 2.0) are cited and used under their licenses. The words and answer keys of published NYT Connections puzzles are used for research evaluation; the repository documents their provenance, and NYT retains all rights in the underlying puzzles.

13. New Assets

Question: Are new assets introduced in the paper well documented, and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released code and configs are documented in the repository README; released adapters and result sets carry documentation cards stating training provenance and research-only intended use, including the disclosed prompt-order bug.

14. Crowdsourcing and Research with Human Subjects

Question: Does the paper include the full text of instructions given to participants, screen-shots, and details about compensation?

Answer: [N/A]

Justification: No crowdsourcing and no human subjects.

15. IRB Approvals

Question: Does the paper describe potential risks to study participants and whether IRB approvals were obtained?

Answer: [N/A]

Justification: No human subjects.

16. Declaration of LLM Usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods?

Answer: [Yes]

Justification: LLM judges are a declared component of the method: §8 names the three models (gpt-5-mini, gemini-2.5-flash, claude-sonnet-4-5), the blind scoring protocol, and their correlated-scorer limitation.